

A Brief Introduction to C++

Overview

- C++ Programming Model
- Coding Conventions (used in COS30008)
- A Simple Particle Simulation: A First Example
 - Object Construction
 - Operators and Member Functions
 - Class Composition
- Object-oriented Programming in C++

References

- Stanley B. Lippman, Josée Lajoie, and Barbara E. Moo: C++ Primer. 5th Edition. Addison-Wesley (2013)
- Scott Meyers: Effective Modern C++. O'Reilly (2014)
- Anthony Williams: C++ Concurrency in Action - Practical Multithreading. Manning Publications Co. (2012)

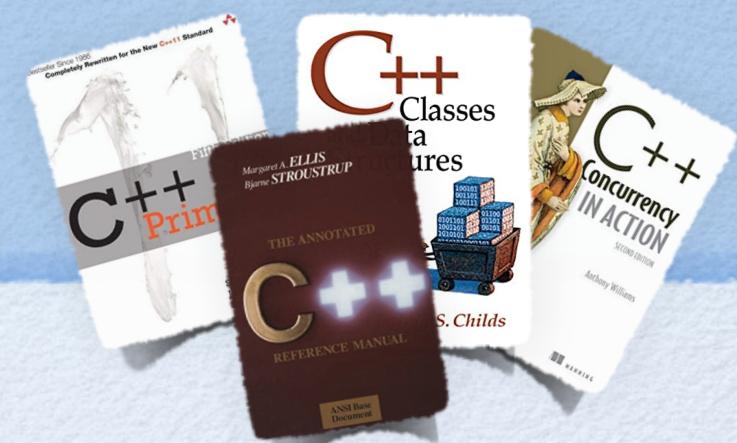
Why C++

- We need to know more than just Java or C#.
- C++ is highly efficient and provides a better match to implement low-level software layers like device controllers or networking protocols.
- C++ is being widely used to develop commercial applications and is at the center of operating system and modern game development.
- Memory is tangible in C++ and we can, therefore, study the effects of design decisions on memory management more directly.

Core Properties of Programming Languages

- Programming languages provide us with a framework to organize computational complexity in our own minds.
- Programming languages offer us the means by which we communicate our understanding about a computerized problem solution.

What is C++



- C++ is a general-purpose, high-level programming language with low-level features.
- Bjarne Stroustrup developed C++ (C with Classes) in 1983 at Bell Labs as an enhancement to the C programming language.
- A first C++ programming language standard was ratified in 1998 as ISO/IEC 14882:1998. The current extending version is ISO/IEC 14882:2011 (informally known as C++11, C++14, C++17, C++20,...).
- C++11 differs greatly from C++98. C++14 differs ...
- C++11 introduces move semantics and lambda expressions (revised in C++14)

Design Philosophy of C++

- C++ is a **hybrid, statically-typed, general-purpose** language that is as efficient and portable as C.
- C++ directly supports **multiple programming styles** like procedural programming, object-oriented programming, or generic programming.
- C++ gives the programmer choice, even if this makes it possible for the programmer to **choose incorrectly!**
- C++ avoids features that are platform specific or not general purpose, but is itself **platform-dependent**.
- C++ does not incur overhead for features that are not used.
- C++ functions without an integrated and sophisticated programming environment.

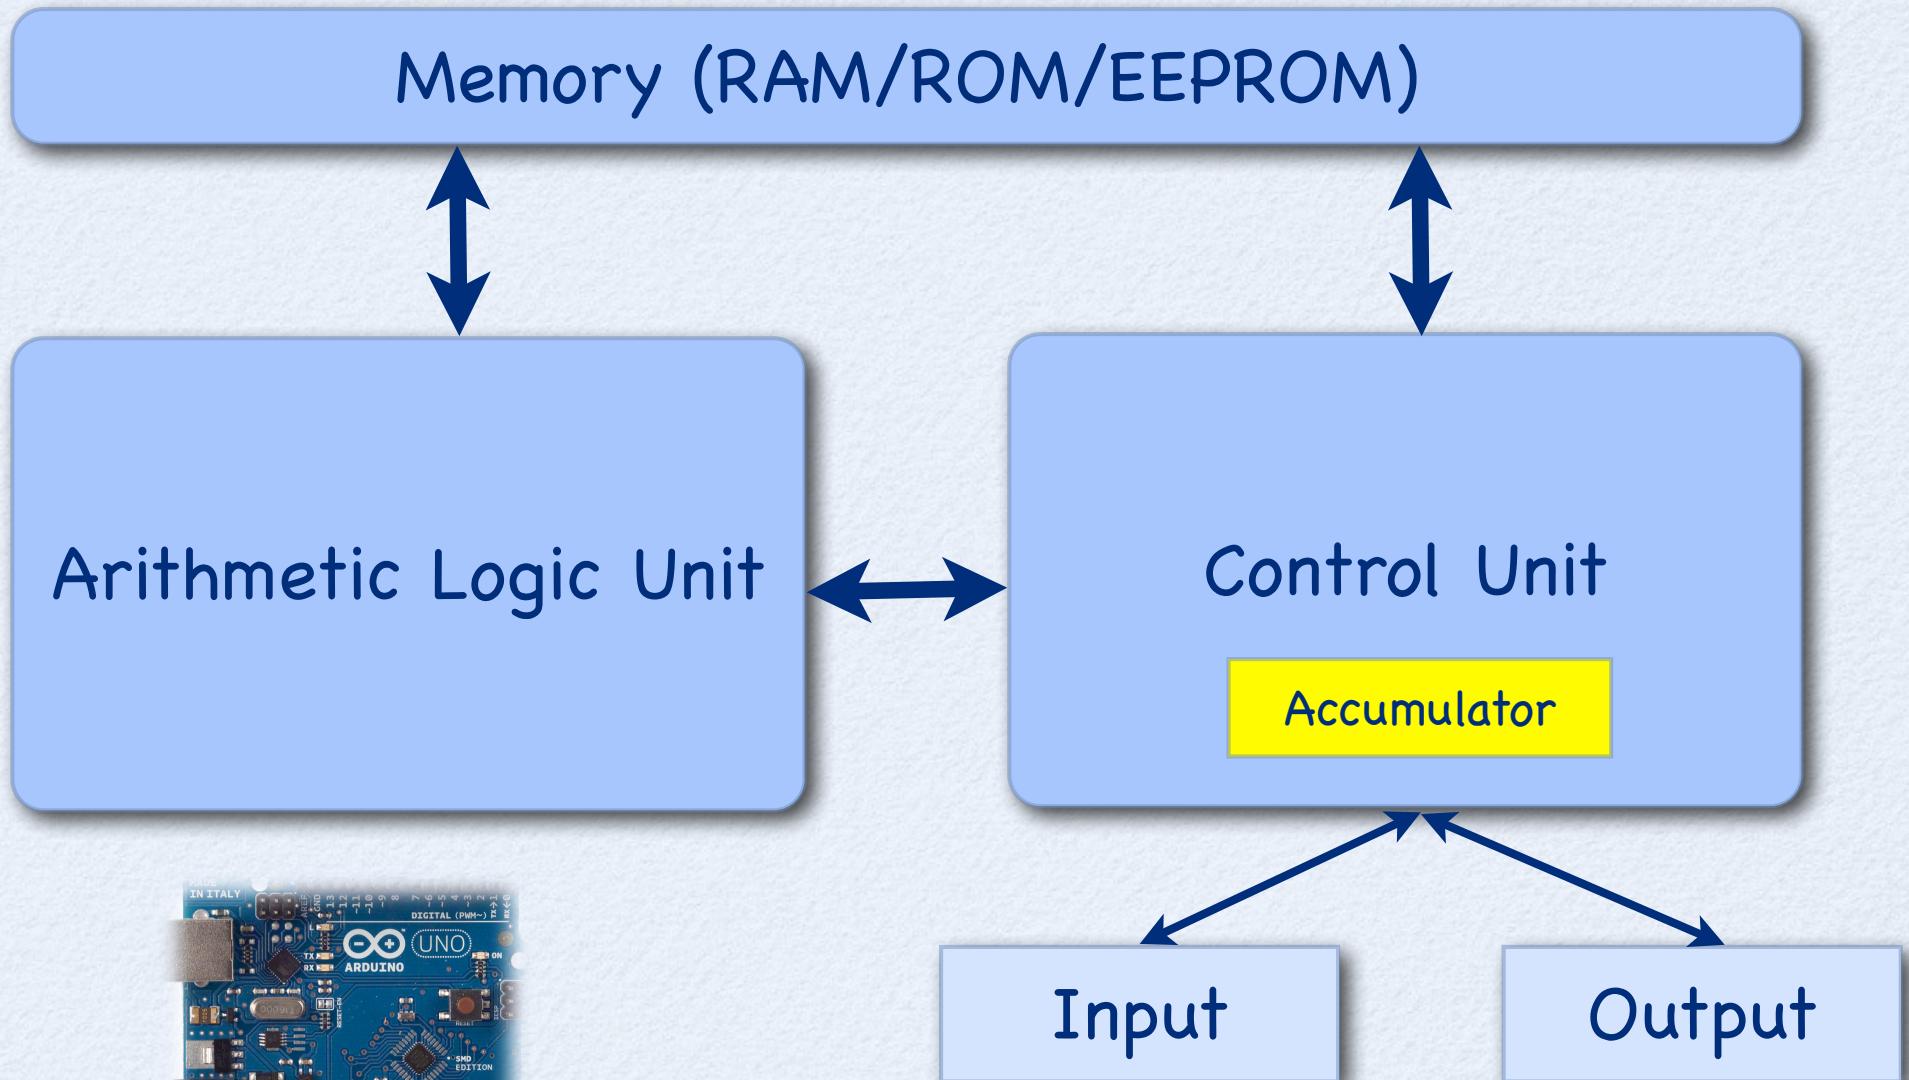
C++ Paradigms

- C++ is a multi-paradigm language.
- C++ provides natural support for
 - the imperative paradigm and
 - the object-oriented paradigm.
- Paradigms must be mixed in any non-trivial project.

Imperative Programming

- This is the oldest style of programming, in which the algorithm for the computation is expressed explicitly in terms of instructions such as assignments, tests, branching and so on.
- Execution of the algorithm requires data values to be held in variables which the program can access and modify.
- Imperative programming corresponds naturally to the earliest, basic and still used model for the architecture of the computer, the von Neumann model.

The von Neumann Architecture

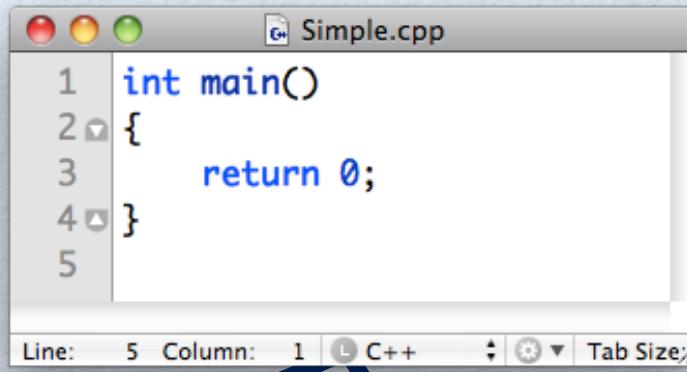


Object-Oriented Programming

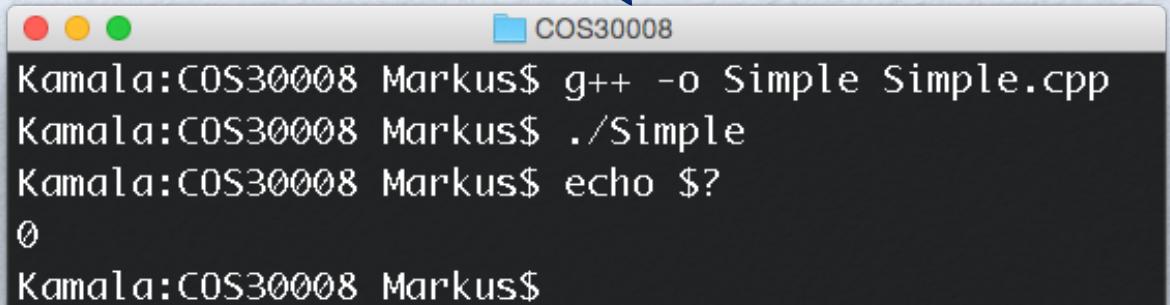
- In general, object-oriented languages are based on the concepts of **class** and **inheritance**, which may be compared to those of **type** and **variable** respectively in a language like Pascal and C.
- A class describes the characteristics common to all its instances, in a form similar to the record of Pascal (structures in C), and thus defines a set of fields.
- In object-oriented programming, instead of applying global procedures or functions to variables, we invoke the **methods** associated with the **instances** (i.e., objects), an action called "**message passing**."
- The basic concept inheritance is used to derive new classes from exiting ones by modifying or extending the inherited class(es).

The Simplest Possible C++ Program

```
.text
    .align 1,0x90
.globl _main
_main:
LFB2:
    pushl %ebp
LCFI0:
    movl %esp, %ebp
LCFI1:
    subl $8, %esp
LCFI2:
    movl $0, %eax
    leave
    ret
LFE2:
    .globl _main.eh
_main.eh = 0
.no_dead_strip _main.eh
.constructor
.destructor
.align 1
.subsections_via_symbols
```

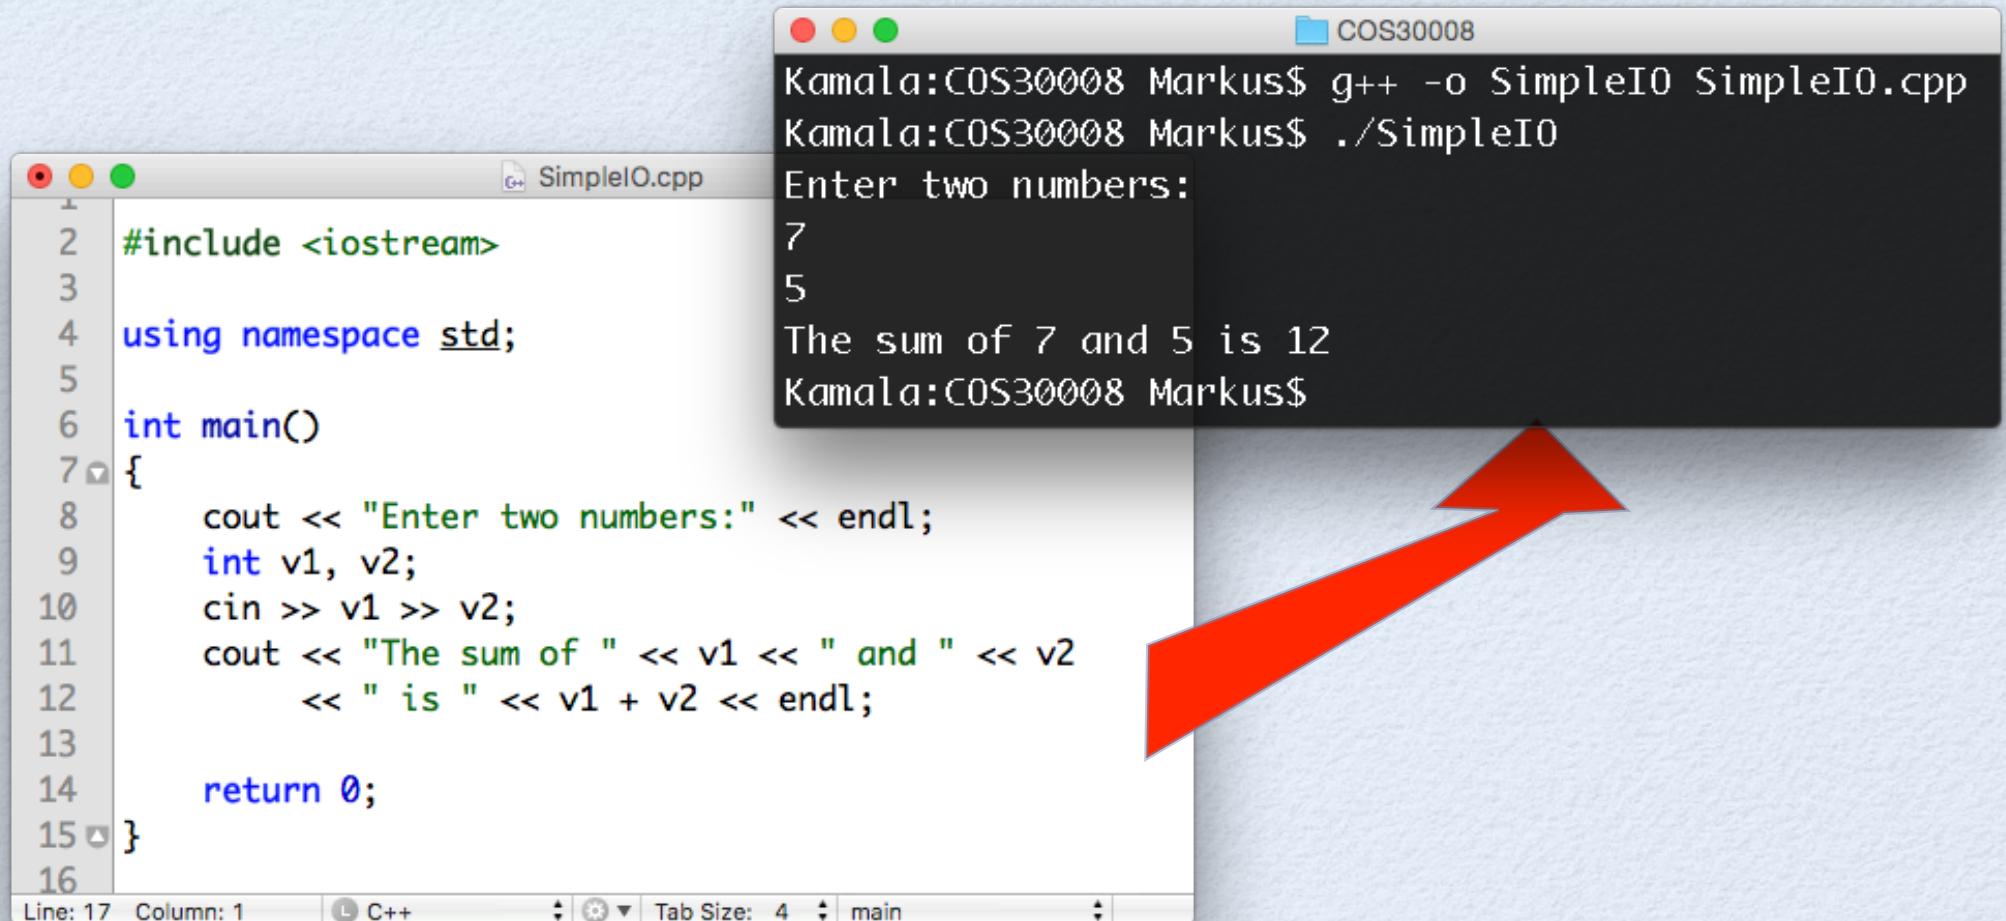


```
Simple.cpp
1 int main()
2 {
3     return 0;
4 }
5
```



```
Line: 5 Column: 1 C++ Tab Size: // COS30008
Kamala:COS30008 Markus$ g++ -o Simple Simple.cpp
Kamala:COS30008 Markus$ ./Simple
Kamala:COS30008 Markus$ echo $?
0
Kamala:COS30008 Markus$
```

Let's make the program more responsive!



A screenshot of a Mac OS X desktop environment. On the left is a code editor window titled "SimpleIO.cpp" containing the following C++ code:

```
1 #include <iostream>
2
3 using namespace std;
4
5 int main()
6 {
7     cout << "Enter two numbers:" << endl;
8     int v1, v2;
9     cin >> v1 >> v2;
10    cout << "The sum of " << v1 << " and " << v2
11        << " is " << v1 + v2 << endl;
12
13    return 0;
14 }
15
16
```

Below the code editor is a terminal window titled "COS30008" with the following output:

```
Kamala:COS30008 Markus$ g++ -o SimpleIO SimpleIO.cpp
Kamala:COS30008 Markus$ ./SimpleIO
Enter two numbers:
7
5
The sum of 7 and 5 is 12
Kamala:COS30008 Markus$
```

A large red arrow points from the text "Let's make the program more responsive!" at the top of the slide down towards the terminal window.

- C++ does not directly define any I/O primitives.
- I/O operations are provided by standard libraries.

The Standard Input Stream `cin`

- `cin` is an object of class `istream` that represents the standard input stream. It corresponds to `stdin` in C.
- `cin` is a globally visible object that is readily available to any C++ compilation unit (i.e., a `.cpp`-file) that includes the library `iostream`.
- `cin` receives input either from the keyboard or a stream associated with the standard input stream.
- We use the operator `>>` to fetch formatted data or use the methods `read` or `get` to retrieve unformatted data from the standard input stream.

The Standard Output Stream cout

- cout is an object of class ostream that represents the standard output stream. It corresponds to stdout in C.
- cout is a globally visible object that is readily available to any C++ compilation unit (i.e., a .cpp-file) that includes the library iostream.
- cout sends data either to the console (as text) or a stream associated with the standard output stream.
- We use the operator << to push formatted data or use the methods write or put to send unformatted data to the standard output stream.

Visual Studio

2019

The screenshot shows the initial opening screen of Microsoft Visual Studio 2022. At the top, the menu bar includes File, Edit, View, Debug, Analyze, Tools, Extensions, Window, Help, and a Search (Ctrl+Q) field. Below the menu is a toolbar with various icons for file operations like Open, Save, and Build. The main area has a large title "What would you like to do?" followed by two sections: "Open recent" and "Get started".

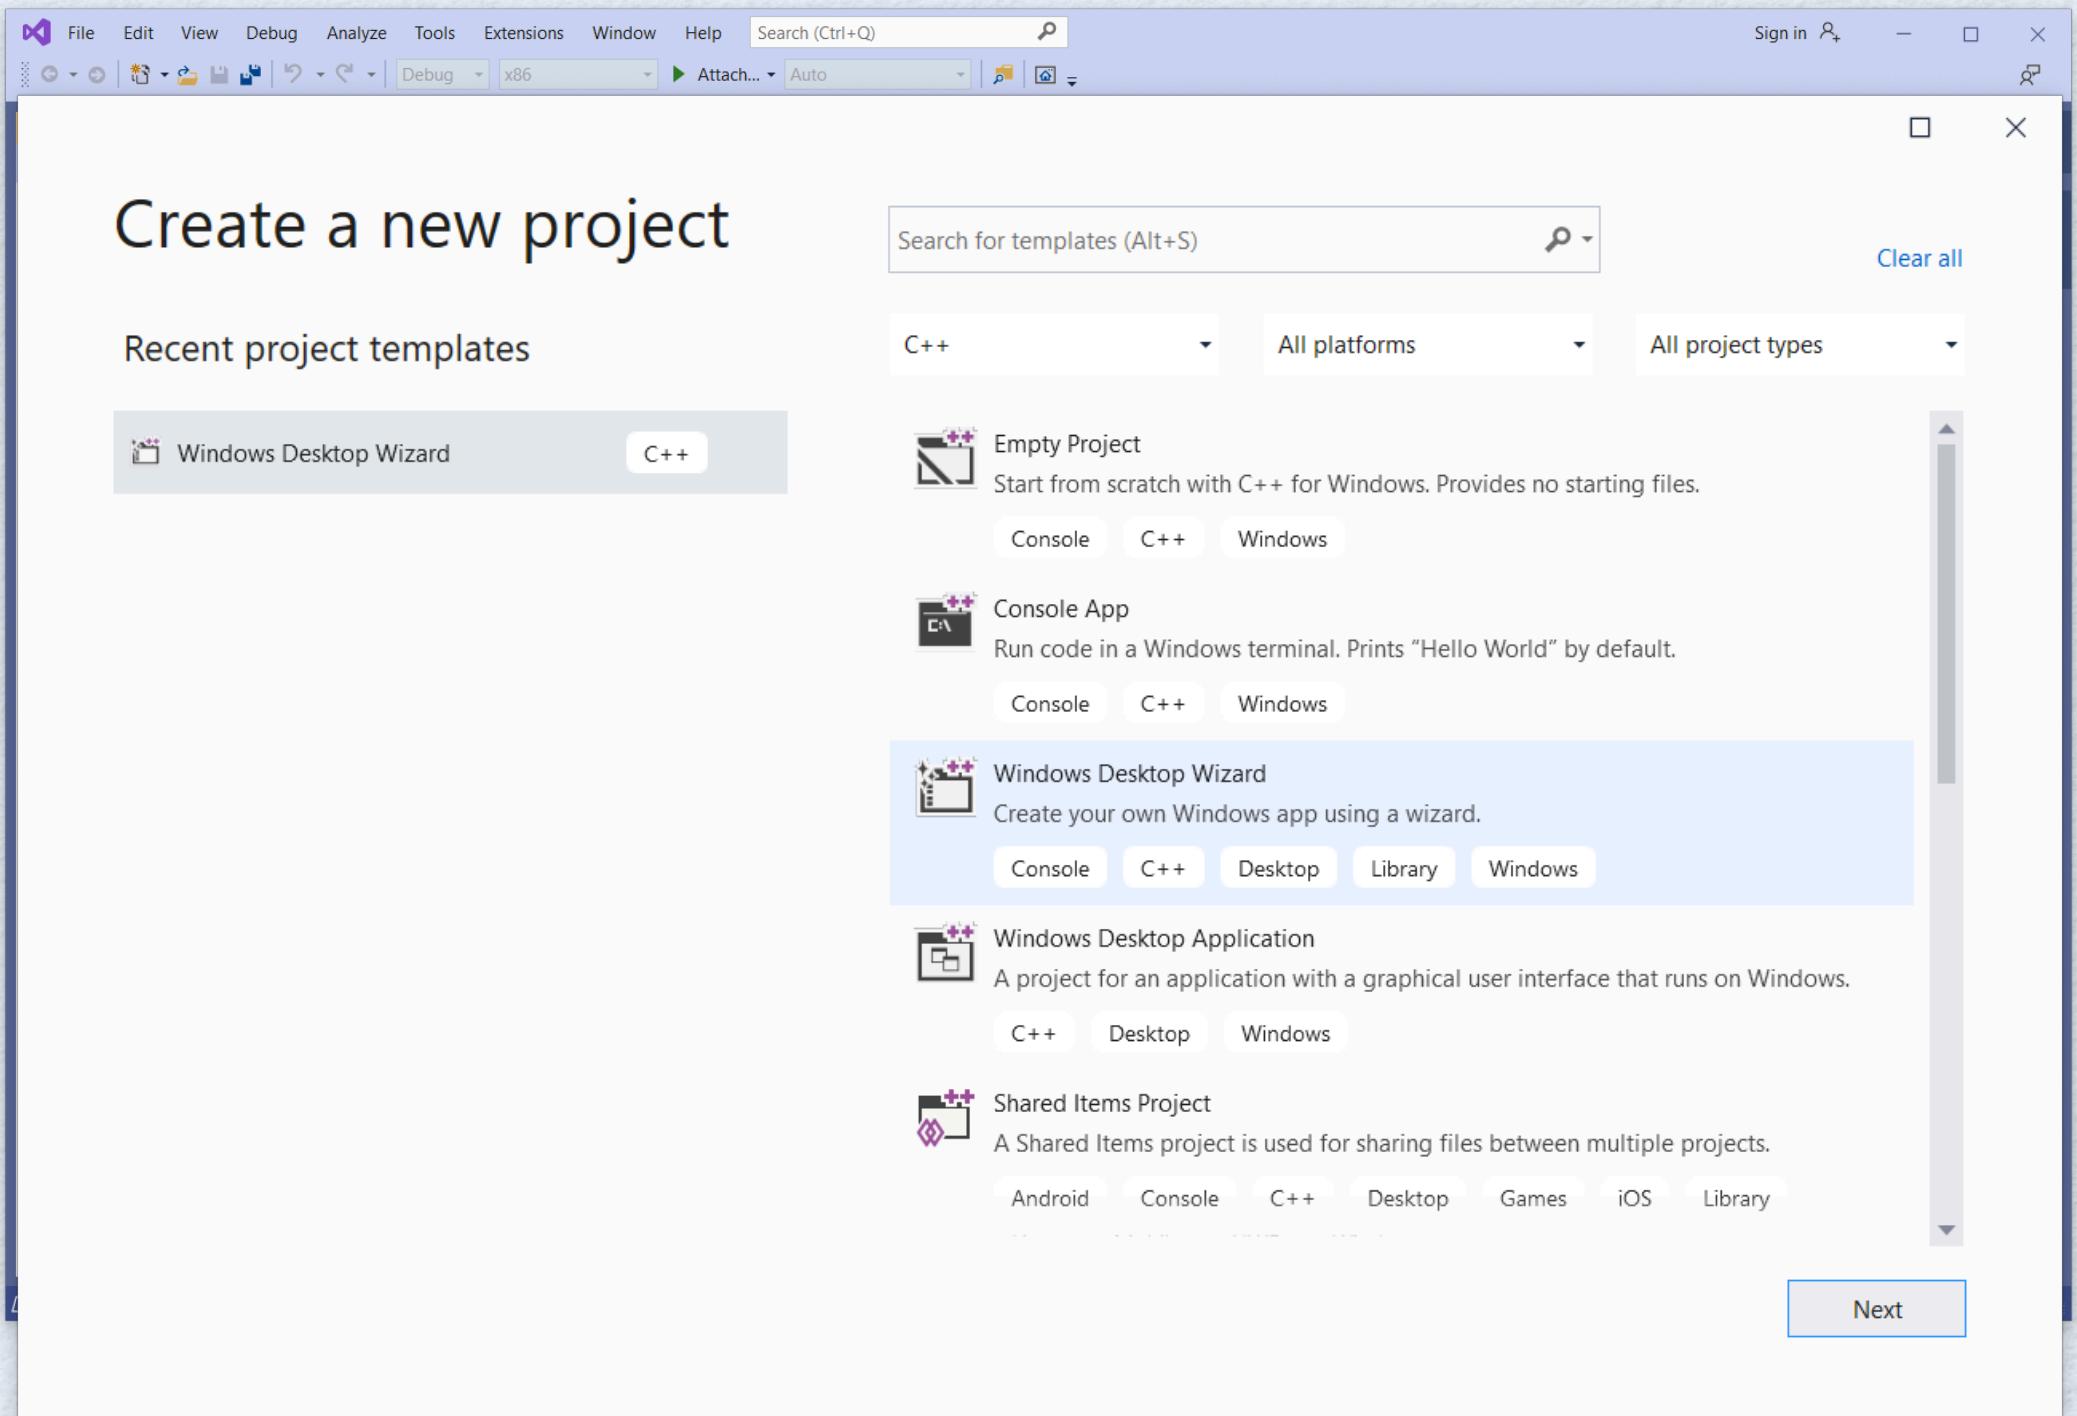
Open recent

As you use Visual Studio, any projects, folders, or files that you open will show up here for quick access.

You can pin anything that you open frequently so that it's always at the top of the list.

Get started

- Clone a repository**
Get code from an online repository like GitHub or Azure DevOps
- Open a project or solution**
Open a local Visual Studio project or .sln file
- Open a local folder**
Navigate and edit code within any folder
- Create a new project**
Choose a project template with code scaffolding to get started



Configure

Windows Desktop

Project name

Location

Solution name i

Place solution and project in same folder

Windows Desktop Project

Application type

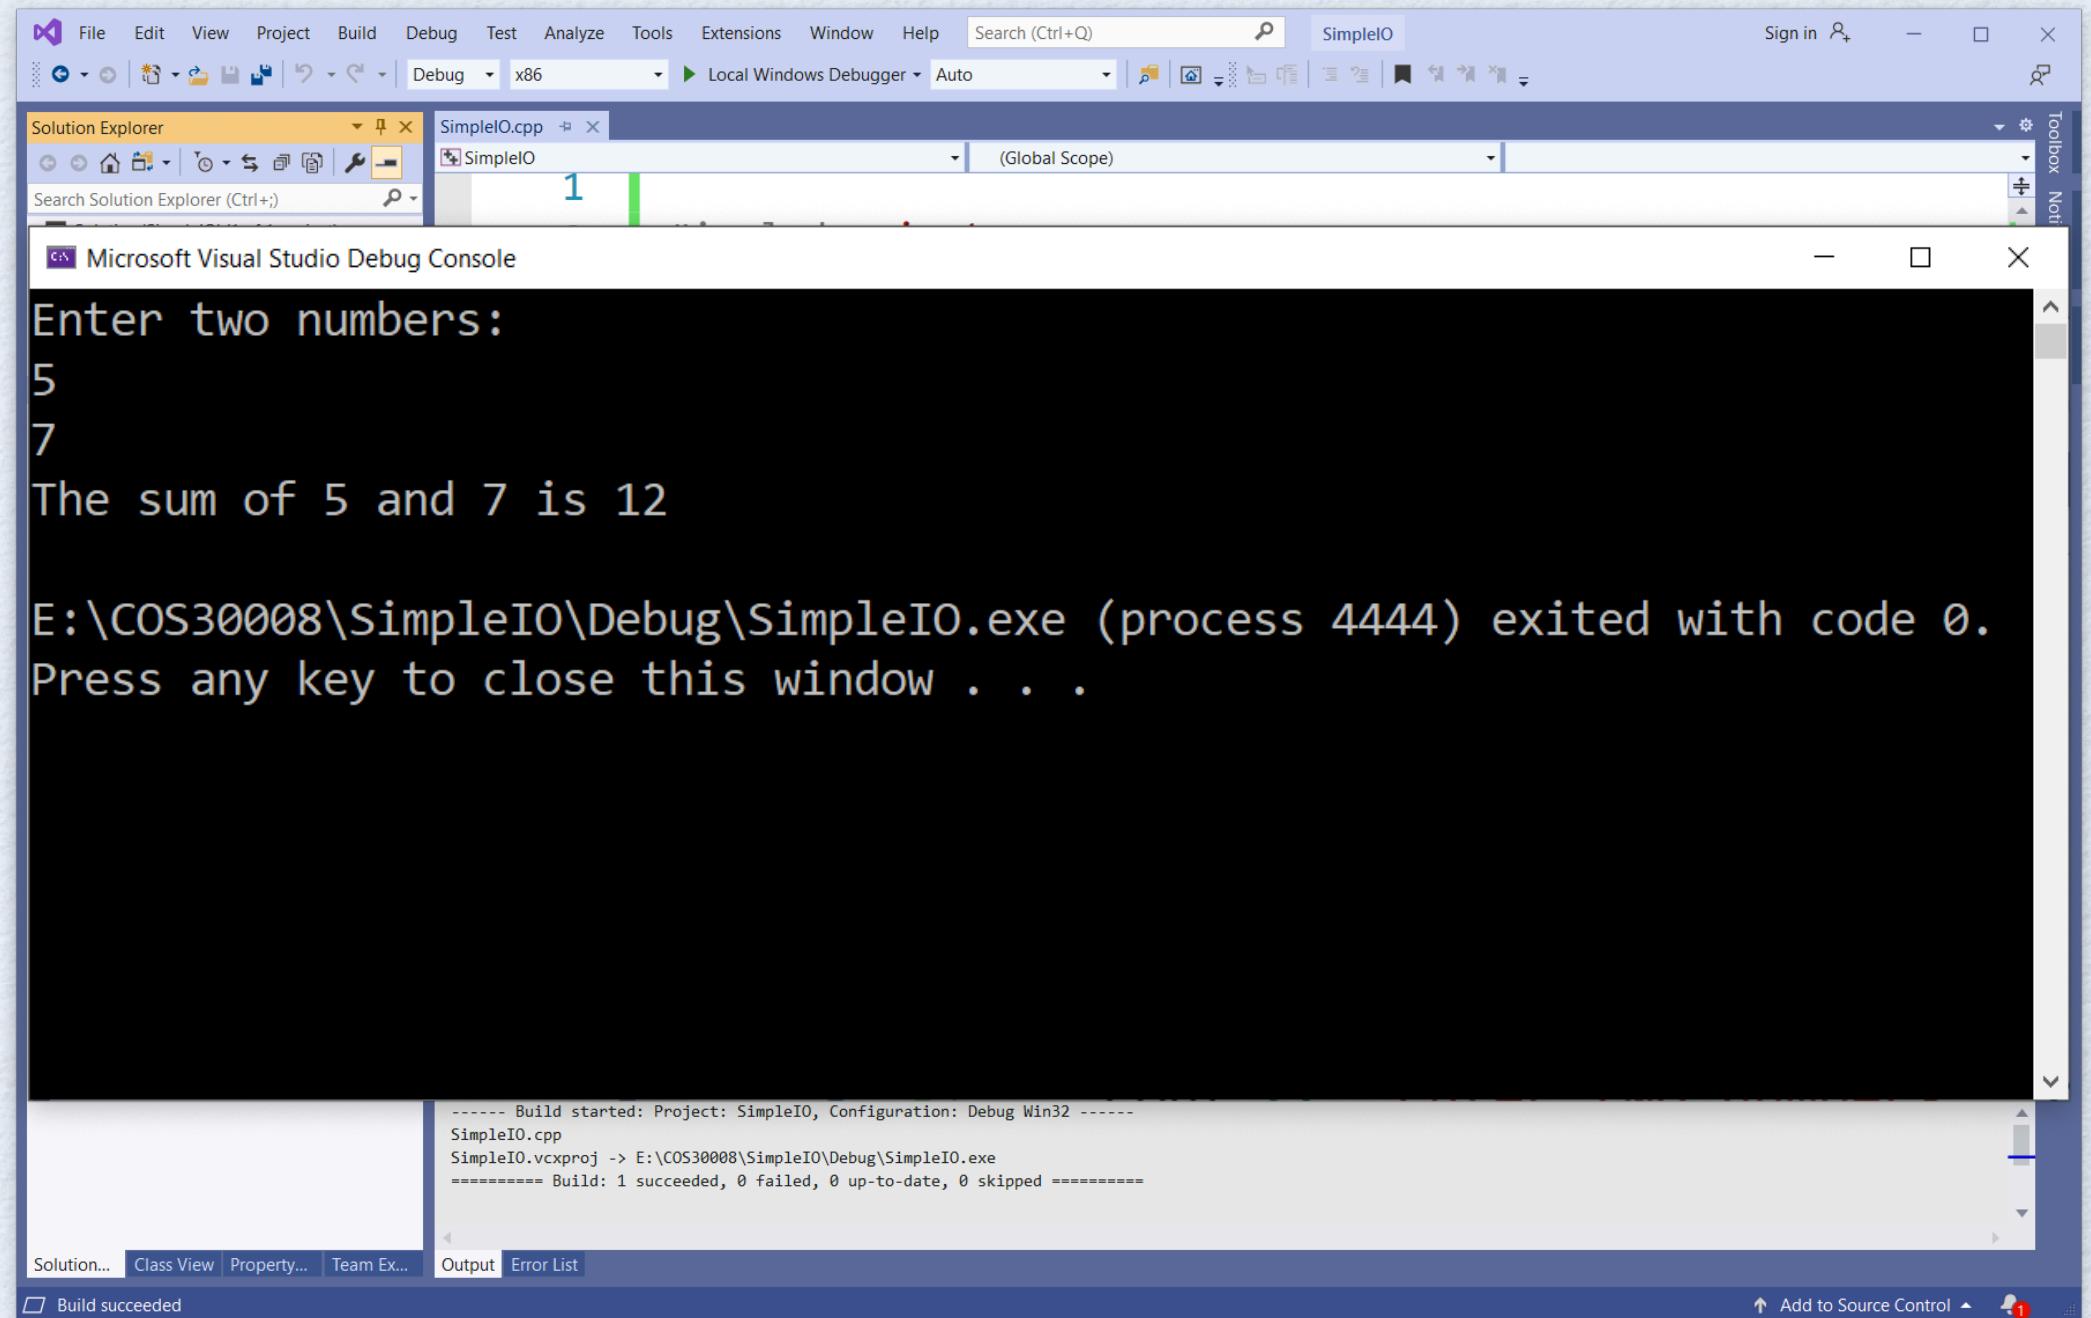
Additional options:

 Empty project Precompiled header Export symbols MFC headers

i Tip: You can also use the Empty Project template to create this kind of project.

The screenshot shows the Microsoft Visual Studio IDE interface. The main window displays a code editor with the file `SimpleIO.cpp` open. The code implements a simple program that prompts the user to enter two numbers, calculates their sum, and prints the result. The code is written in C++ using standard input-output streams (`<iostream>`). The code editor features color-coded syntax highlighting, including red for strings, blue for keywords like `#include` and `int`, and green for comments. The Solution Explorer on the left shows a single project named `SimpleIO` with one source file, `SimpleIO.cpp`. The Output and Error List tabs at the bottom indicate that no issues were found in the code.

```
1 #include <iostream>
2
3 using namespace std;
4
5
6 int main()
7 {
8     cout << "Enter two numbers:" << endl;
9     int v1, v2;
10    cin >> v1 >> v2;
11    cout << "The sum of " << v1 << " and " << v2
12        << " is " << v1 + v2 << endl;
13
14    return 0;
15
16 }
```



A Simple Particle Simulation

- Before we can write a program in a new language, we need to know some of its basic features. C++ is no exception.
- Let's start with a simple particle simulation (based on ideas presented on www.codingmath.com).
- The simulation program requires us to
 - define variables
 - perform input and output
 - define two data structures (i.e., classes) to hold the data we are managing: Vector2D and Particle2D
 - implement basic vector math operations for Vector2D and Particle2D
 - write control code to simulate a particle object

Technical References

- Eric Lengyel: Mathematics for 3D Game Programming and Computer Graphics, Course Technology (2012)
- Stanley B. Lippman, Josée Lajoie, and Barbara E. Moo: C++ Primer. 5th Edition. Addison-Wesley (2013)
- Kenneth H. Rosen: Discrete Mathematics and Its Application. 7th Edition. McGraw-Hill (2012)
- Thomas H. Cormen, Charles E. Leiserson, Ronald L. Rivest, and Clifford Stein: Introduction to Algorithms. 3rd Edition. The MIT Press (2009)
- Paul Orland: Math for Programmers – 3D graphics, machine learning, and simulations with Python. Manning Publications (2020)
- www.codingmath.com

2D Vector Operations

Length:

$$\|\vec{v}\| = \sqrt{x^2 + y^2}$$

Dot Product:

$$\vec{v} \cdot \vec{u} = v_x u_x + v_y u_v = \|\vec{v}\| \|\vec{u}\| \cos \alpha$$

Cross Product*:

$$\vec{v} \times \vec{u} = \det \begin{pmatrix} v_x & u_x \\ v_y & u_y \end{pmatrix} = v_x u_y - u_x v_y$$

Align Orientation:

$$\vec{v}' = \begin{bmatrix} \cos \alpha \\ \sin \alpha \end{bmatrix} \|\vec{v}\|$$

Orientation/Direction:

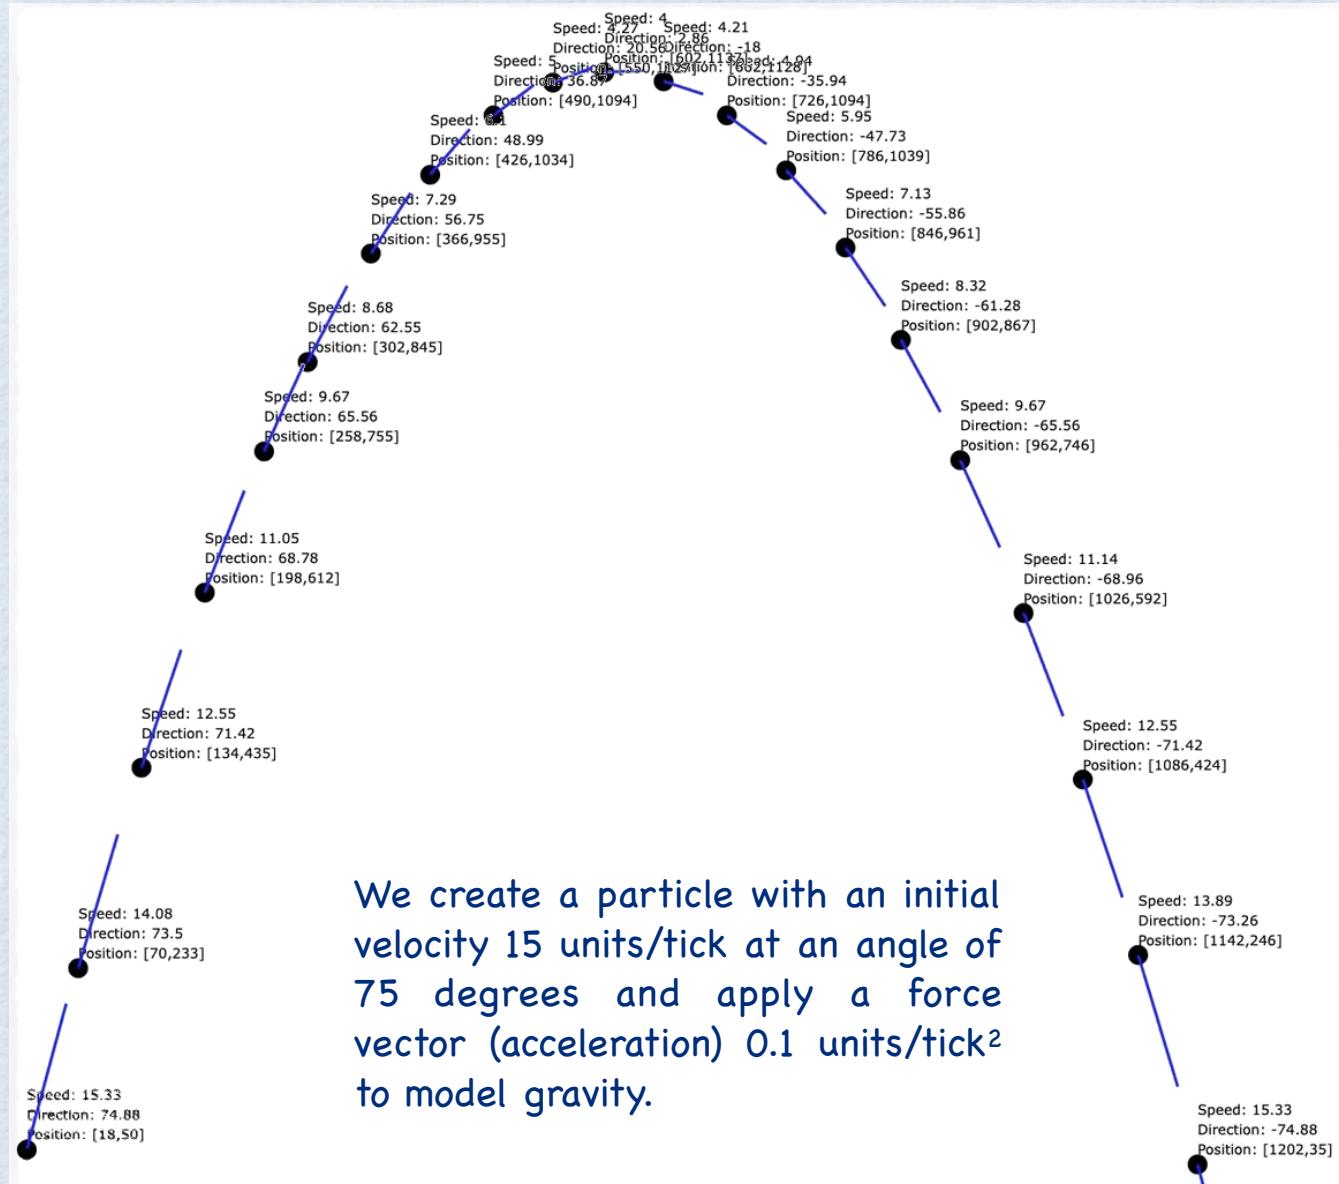
$$\theta = \tan^{-1}(v_y/v_x)$$

Normalization:

$$\vec{n}_v = \frac{\vec{v}}{\|\vec{v}\|}, \|\vec{n}_v\| = 1$$

* The cross product is a three-dimensional concept. In 2D, we use it to determine whether consecutive line segments turn left or right.

Visualization of the Particle Simulation



Preliminaries

Coding Conventions

Coding Conventions

- Coding conventions establish guidelines for a specific programming language that recommend programming style, practices, and methods for each aspect of a program written in this language.
- Coding conventions are only applicable to the human maintainers and peer reviewers of a software project.
- Conventions may be formalized in a documented set of rules that an entire team or company follows, or may be as informal as the habitual coding practices of an individual.
- Coding conventions are not enforced by compilers.

CamelCase

- CamelCase (or medial capitals) is a practice of writing names of variable or functions with some inner uppercase letters to denote embedded words:
 - `InputStreamReader`: character input stream
 - `getEncoding`: getter method for data encoding
 - `MyIntegerArray`: array variable
 - `CreateWindowEx`: Windows function
- Two popular variants:
 - Pascal InfixCaps - the first letter should be a capital, and any embedded words in an identifier should be in caps, as well as any acronym that is embedded.
 - Java - variables are mixed case with a lowercase first letter, methods should be verbs, in mixed case with the first letter lowercase, and classes should be nouns, in mixed case with the first letter of each internal word capitalized.

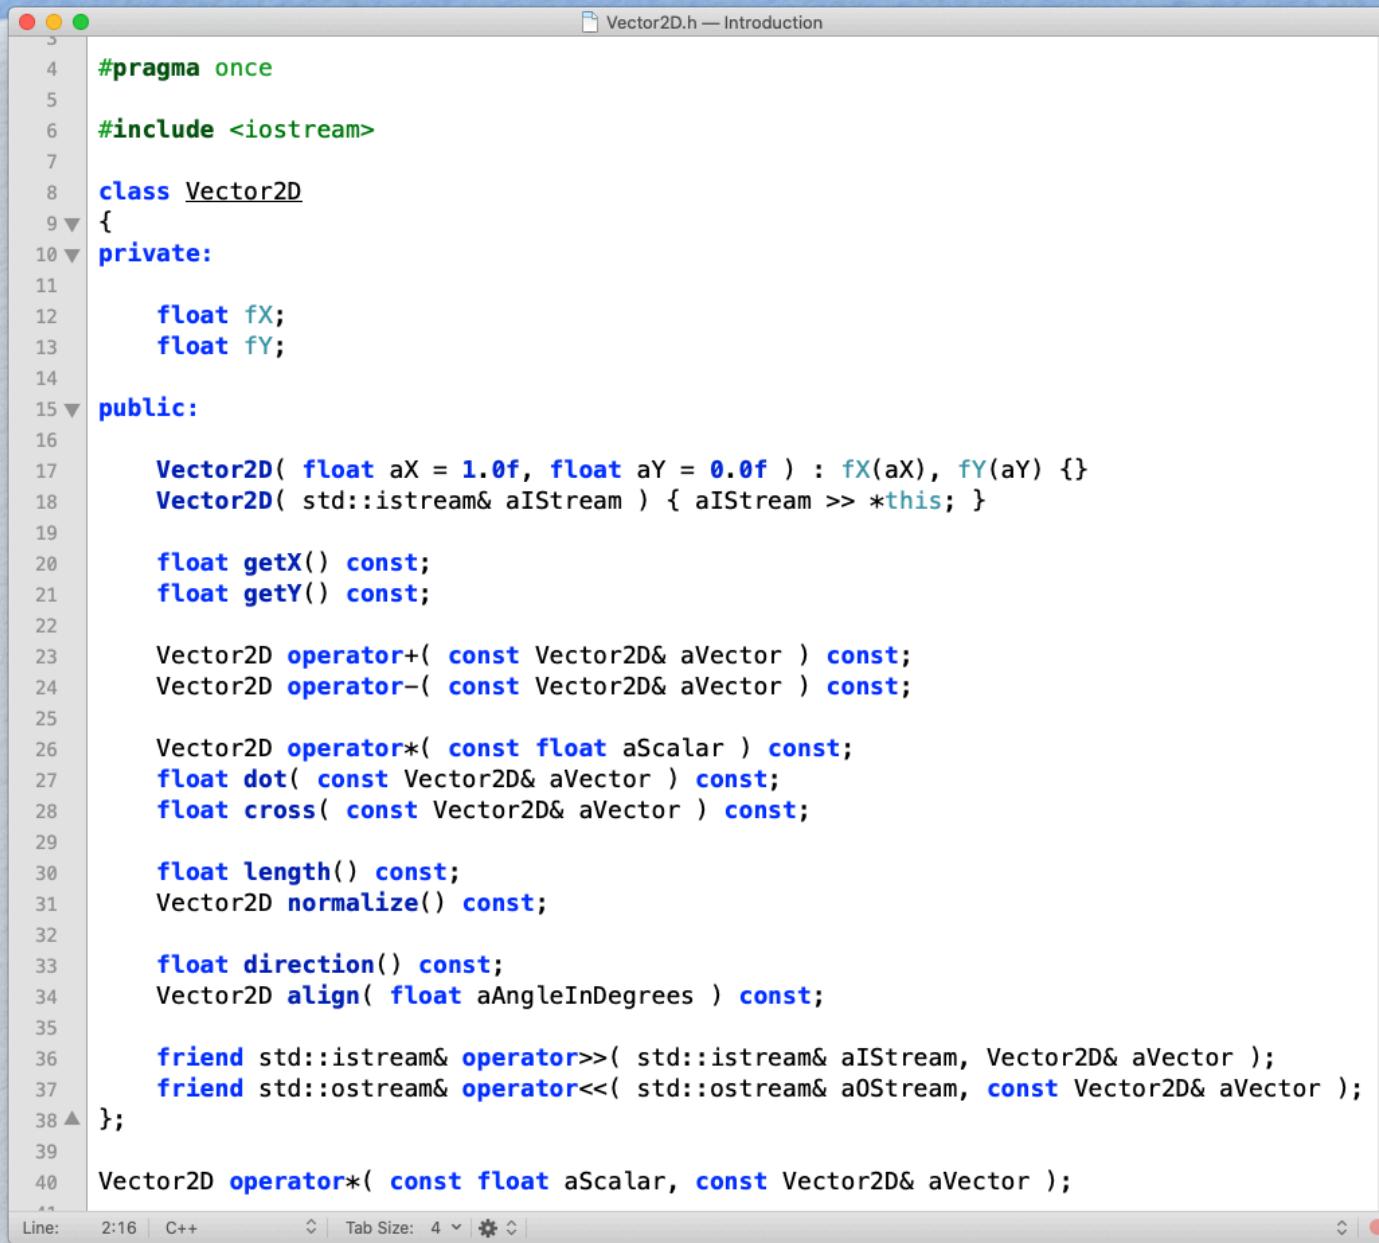
Borland-Inspired Style Guide Elements

- Field names should start with the letter 'f'.
- Local variable names should start with the letter 'l'.
- Parameter names should start with the letter 'a'.
- Global variable names should start with the letter 'g'.

Which Coding Standard To Use

- Coding conventions are not enforced by compilers.
- Not following some or all of the rules has no impact on the executable programs created from the source code.
- Code standards facilitate program comprehension and make program maintenance easier.
- Every organization may enforce some sort of coding standards as part of organization's quality assurance process.

Vector2D: A Basic 2D Vector Class



The screenshot shows a code editor window titled "Vector2D.h — Introduction". The code is written in C++ and defines a class named Vector2D. The class has private members fX and fY, and public methods for construction, reading from an istream, getting coordinates, performing arithmetic operations (addition, subtraction, multiplication by scalar), calculating length and direction, normalizing, aligning, and writing to an ostream.

```
#pragma once
#include <iostream>
class Vector2D
{
private:
    float fX;
    float fY;
public:
    Vector2D( float aX = 1.0f, float aY = 0.0f ) : fX(aX), fY(aY) {}
    Vector2D( std::istream& aInputStream ) { aInputStream >> *this; }

    float getX() const;
    float getY() const;

    Vector2D operator+( const Vector2D& aVector ) const;
    Vector2D operator-( const Vector2D& aVector ) const;

    Vector2D operator*( const float aScalar ) const;
    float dot( const Vector2D& aVector ) const;
    float cross( const Vector2D& aVector ) const;

    float length() const;
    Vector2D normalize() const;

    float direction() const;
    Vector2D align( float aAngleInDegrees ) const;

    friend std::istream& operator>>( std::istream& aInputStream, Vector2D& aVector );
    friend std::ostream& operator<<( std::ostream& aOutputStream, const Vector2D& aVector );
};

Vector2D operator*( const float aScalar, const Vector2D& aVector );
```

The Structure of a Class

```
class X
{
    private:
        // private members
    protected:
        // protected members
    public:
        // public members
};
```

Don't forget the semicolon!

Access Modifiers

- **public:**
 - Public members can be accessed anywhere, including outside of the class itself.
- **protected:**
 - Protected members can be accessed within the class in which they are declared and within derived classes.
- **private:**
 - Private members can be accessed only within the class in which they are declared.

Vector2D: Private Members



```
#pragma once
#include <iostream>
class Vector2D
{
private:
    float fX;
    float fY;

public:
    Vector2D( float aX = 1.0f, float aY = 0.0f ) : fX(aX), fY(aY) {}
    Vector2D( std::istream& aInputStream ) { aInputStream >> *this; }

    float getX() const;
    float getY() const;

    Vector2D operator+( const Vector2D& aVector ) const;
    Vector2D operator-( const Vector2D& aVector ) const;

    Vector2D operator*( const float aScalar ) const;
    float dot( const Vector2D& aVector ) const;
    float cross( const Vector2D& aVector ) const;

    float length() const;
    Vector2D normalize() const;

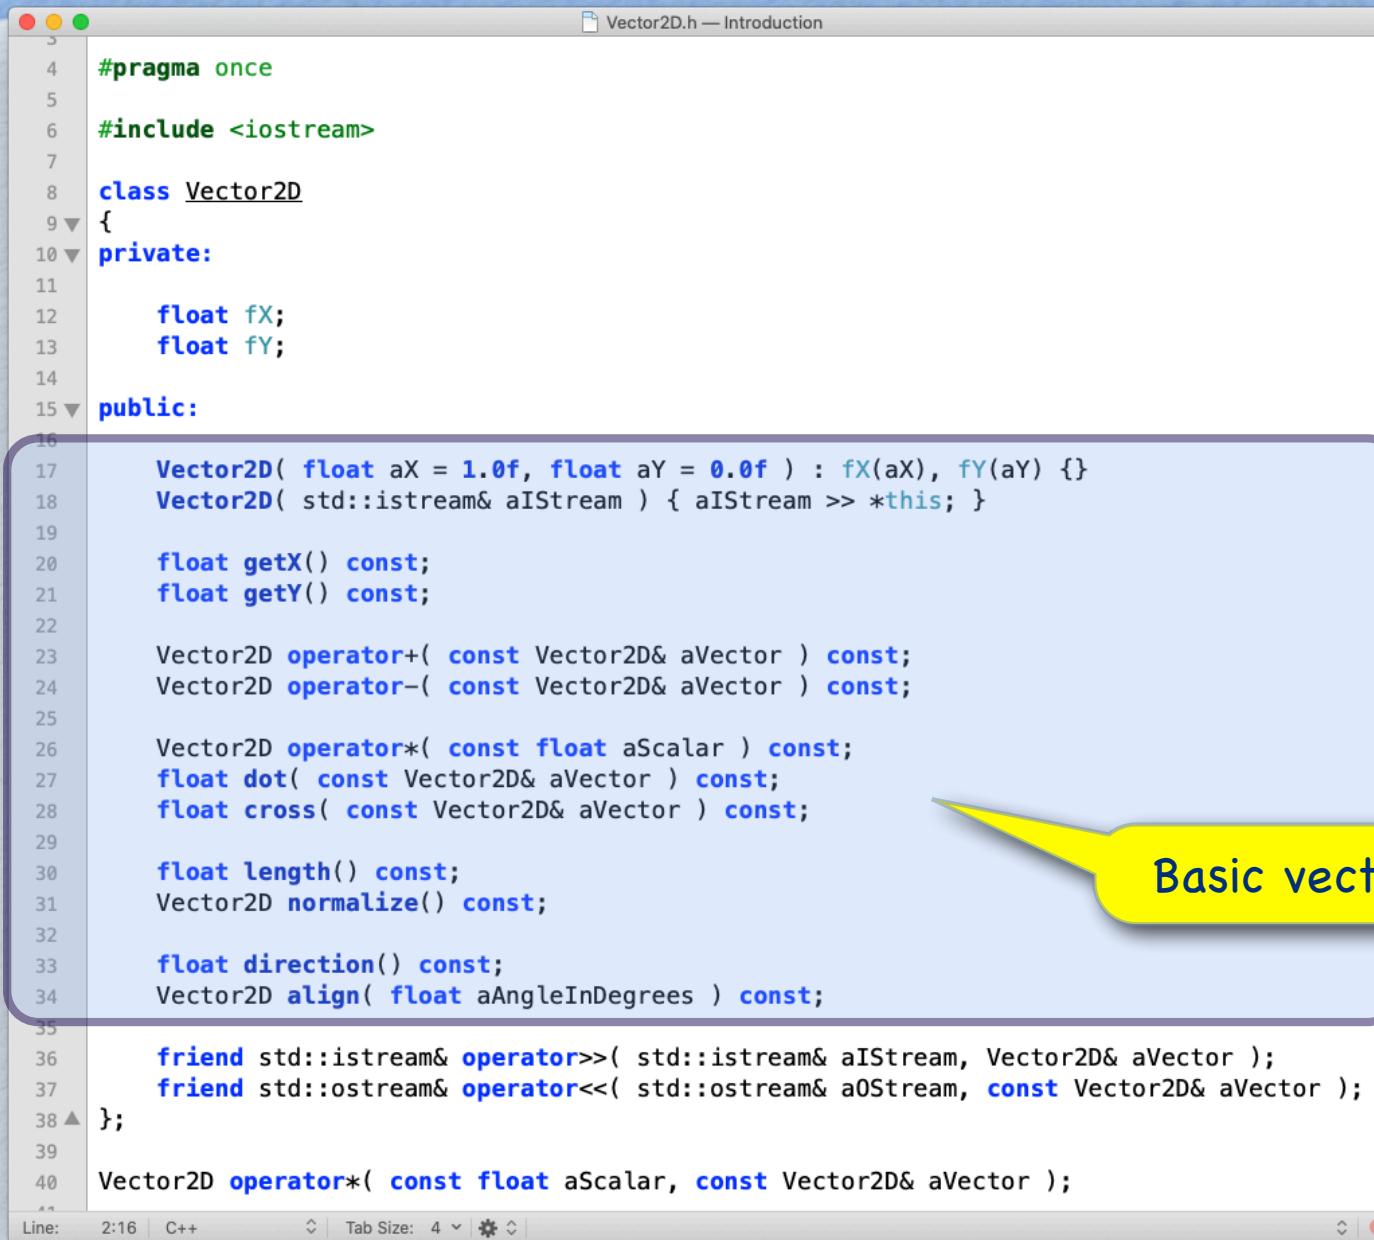
    float direction() const;
    Vector2D align( float aAngleInDegrees ) const;

    friend std::istream& operator>>( std::istream& aInputStream, Vector2D& aVector );
    friend std::ostream& operator<<( std::ostream& aOutputStream, const Vector2D& aVector );
};

Vector2D operator*( const float aScalar, const Vector2D& aVector );
```

Line: 2:16 | C++ | Tab Size: 4 | 

Vector2D: Public Members



```
#pragma once
#include <iostream>
class Vector2D
{
private:
    float fX;
    float fY;
public:
    Vector2D( float aX = 1.0f, float aY = 0.0f ) : fX(aX), fY(aY) {}
    Vector2D( std::istream& aInputStream ) { aInputStream >> *this; }

    float getX() const;
    float getY() const;

    Vector2D operator+( const Vector2D& aVector ) const;
    Vector2D operator-( const Vector2D& aVector ) const;

    Vector2D operator*( const float aScalar ) const;
    float dot( const Vector2D& aVector ) const;
    float cross( const Vector2D& aVector ) const;

    float length() const;
    Vector2D normalize() const;

    float direction() const;
    Vector2D align( float aAngleInDegrees ) const;

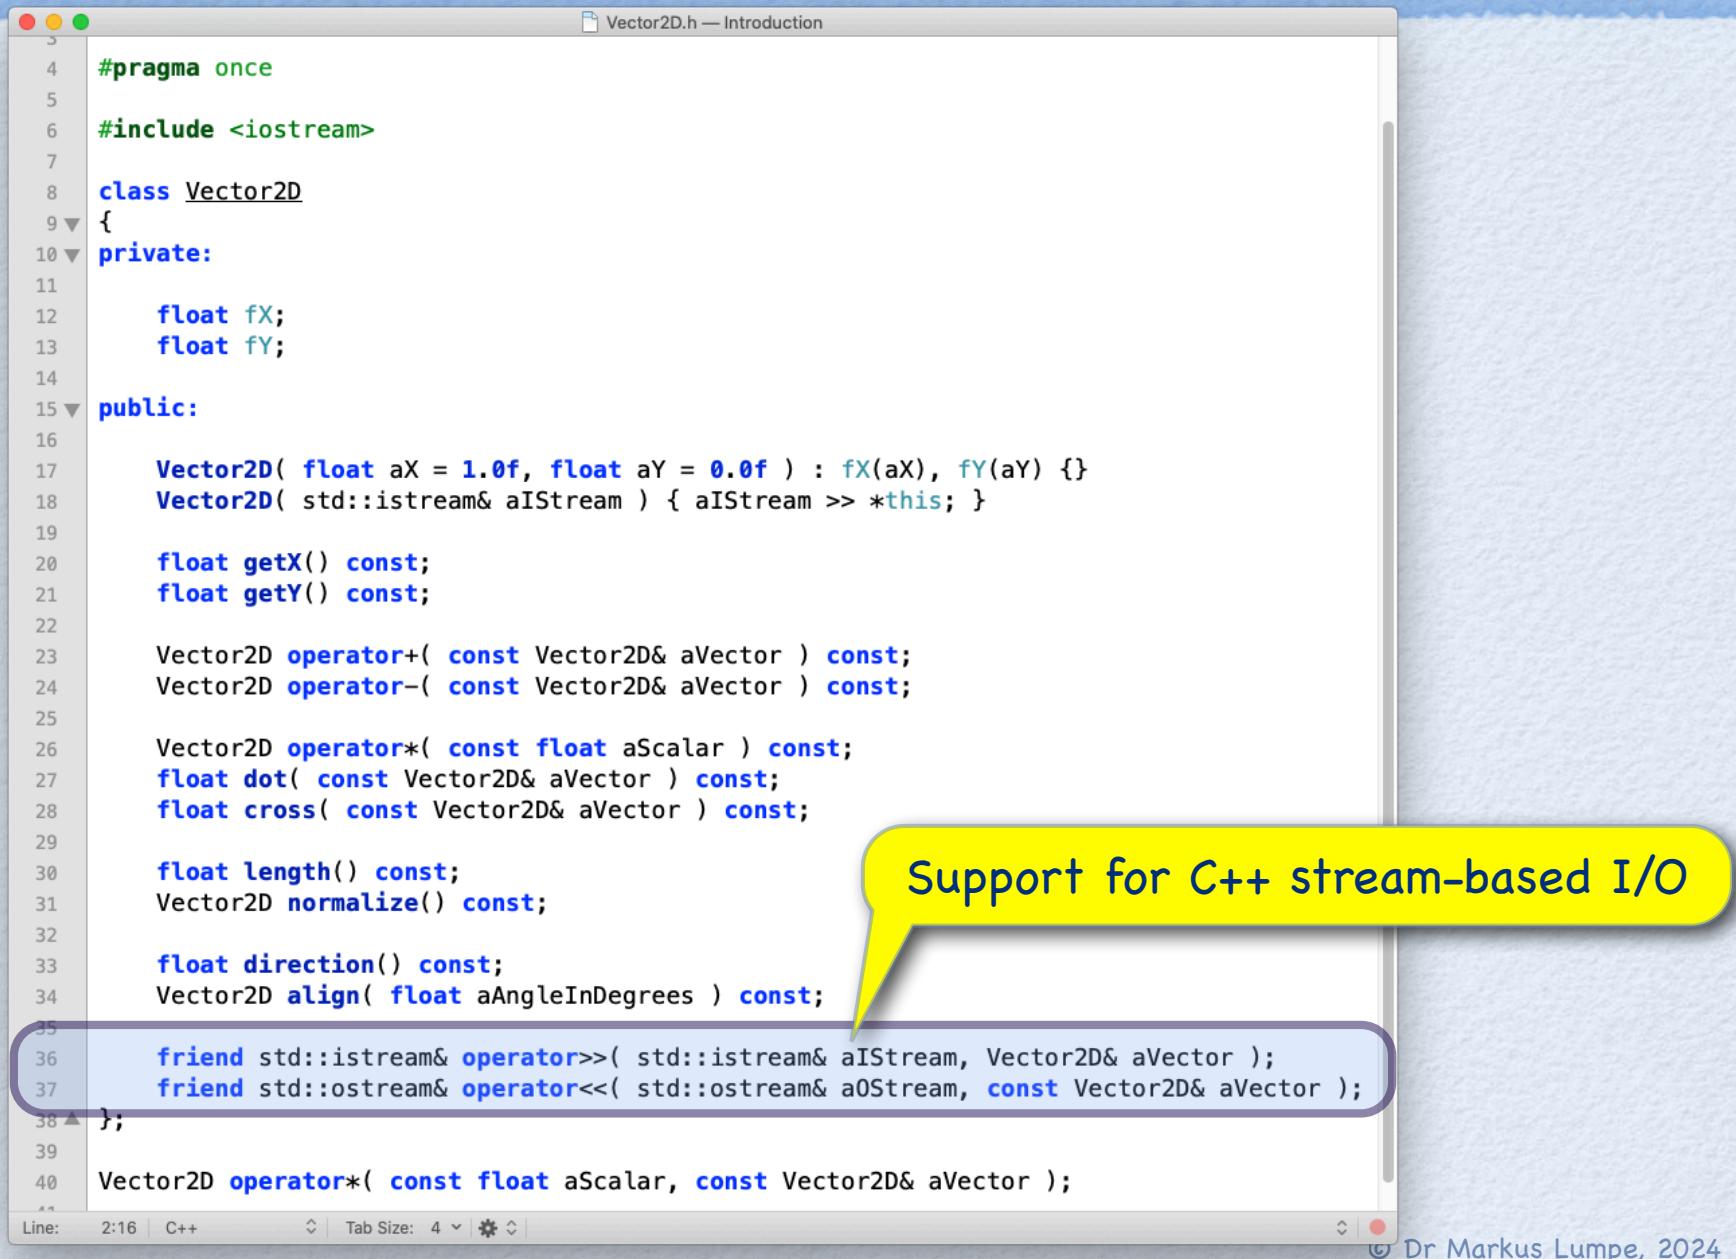
    friend std::istream& operator>>( std::istream& aInputStream, Vector2D& aVector );
    friend std::ostream& operator<<( std::ostream& aOutputStream, const Vector2D& aVector );
};

Vector2D operator*( const float aScalar, const Vector2D& aVector );
```

Line: 2:16 | C++ | Tab Size: 4 | 39 | Dr Markus Lumpe, 2024

Basic vector operations

Vector2D: Friends



```
#pragma once
#include <iostream>
class Vector2D
{
private:
    float fX;
    float fY;
public:
    Vector2D( float aX = 1.0f, float aY = 0.0f ) : fX(aX), fY(aY) {}
    Vector2D( std::istream& aInputStream ) { aInputStream >> *this; }

    float getX() const;
    float getY() const;

    Vector2D operator+( const Vector2D& aVector ) const;
    Vector2D operator-( const Vector2D& aVector ) const;

    Vector2D operator*( const float aScalar ) const;
    float dot( const Vector2D& aVector ) const;
    float cross( const Vector2D& aVector ) const;

    float length() const;
    Vector2D normalize() const;

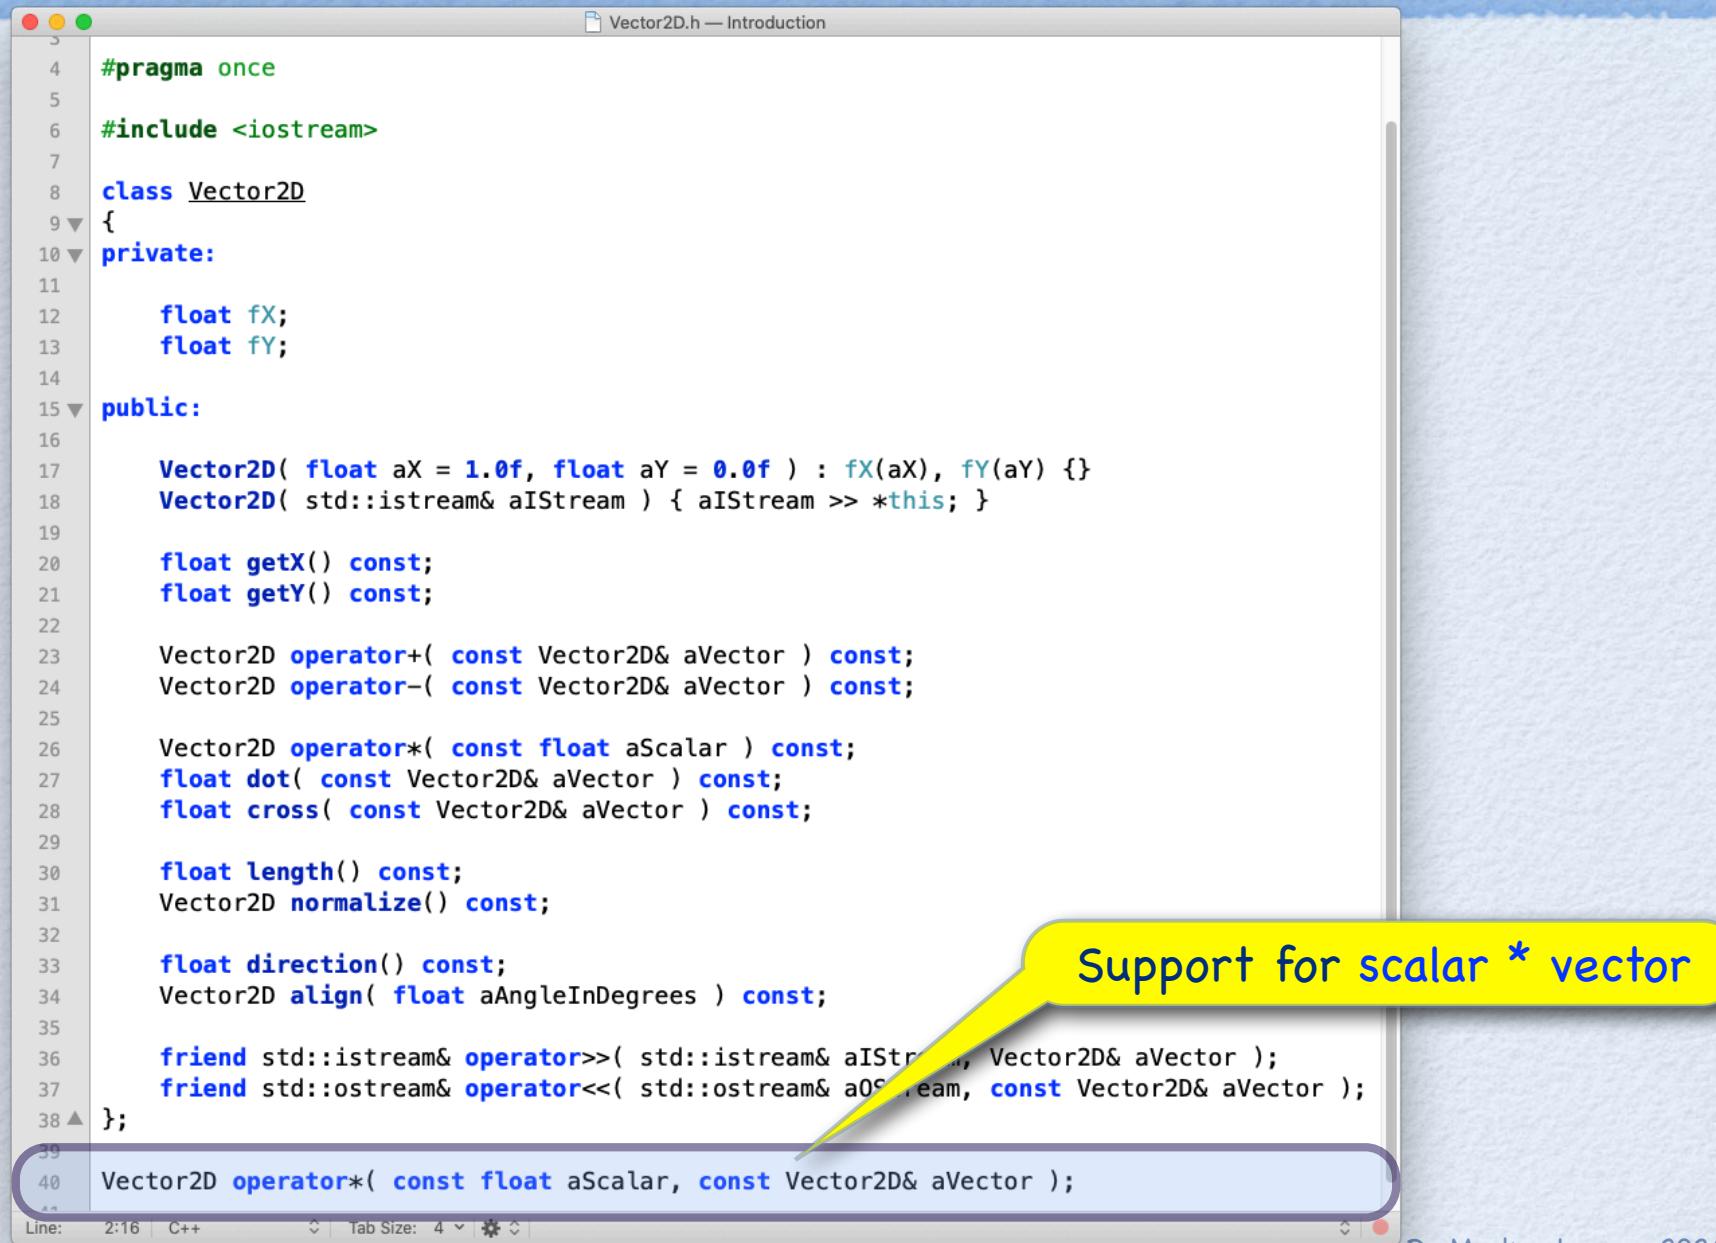
    float direction() const;
    Vector2D align( float aAngleInDegrees ) const;

    friend std::istream& operator>>( std::istream& aInputStream, Vector2D& aVector );
    friend std::ostream& operator<<( std::ostream& aOutputStream, const Vector2D& aVector );
};

Vector2D operator*( const float aScalar, const Vector2D& aVector );
```

Support for C++ stream-based I/O

Vector2D: Ad hoc Definitions



```
#pragma once
#include <iostream>
class Vector2D
{
private:
    float fX;
    float fY;
public:
    Vector2D( float aX = 1.0f, float aY = 0.0f ) : fX(aX), fY(aY) {}
    Vector2D( std::istream& aInputStream ) { aInputStream >> *this; }

    float getX() const;
    float getY() const;

    Vector2D operator+( const Vector2D& aVector ) const;
    Vector2D operator-( const Vector2D& aVector ) const;

    Vector2D operator*( const float aScalar ) const;
    float dot( const Vector2D& aVector ) const;
    float cross( const Vector2D& aVector ) const;

    float length() const;
    Vector2D normalize() const;

    float direction() const;
    Vector2D align( float aAngleInDegrees ) const;

    friend std::istream& operator>>( std::istream& aInputStream, Vector2D& aVector );
    friend std::ostream& operator<<( std::ostream& aOutputStream, const Vector2D& aVector );
};

Vector2D operator*( const float aScalar, const Vector2D& aVector );
```

Support for scalar * vector

const

Constants and Const Qualifier

- What are the problems with

```
for ( int index = 0; index < 128; index++ ) { ... }
```

What is 128?

- We can do better

```
for ( int index = 0; index < BufferSize; index++ ) { ... }
```

Is it safe?

- Defining a const object:

```
const int BufferSize = 128; // initialized at compile time
```

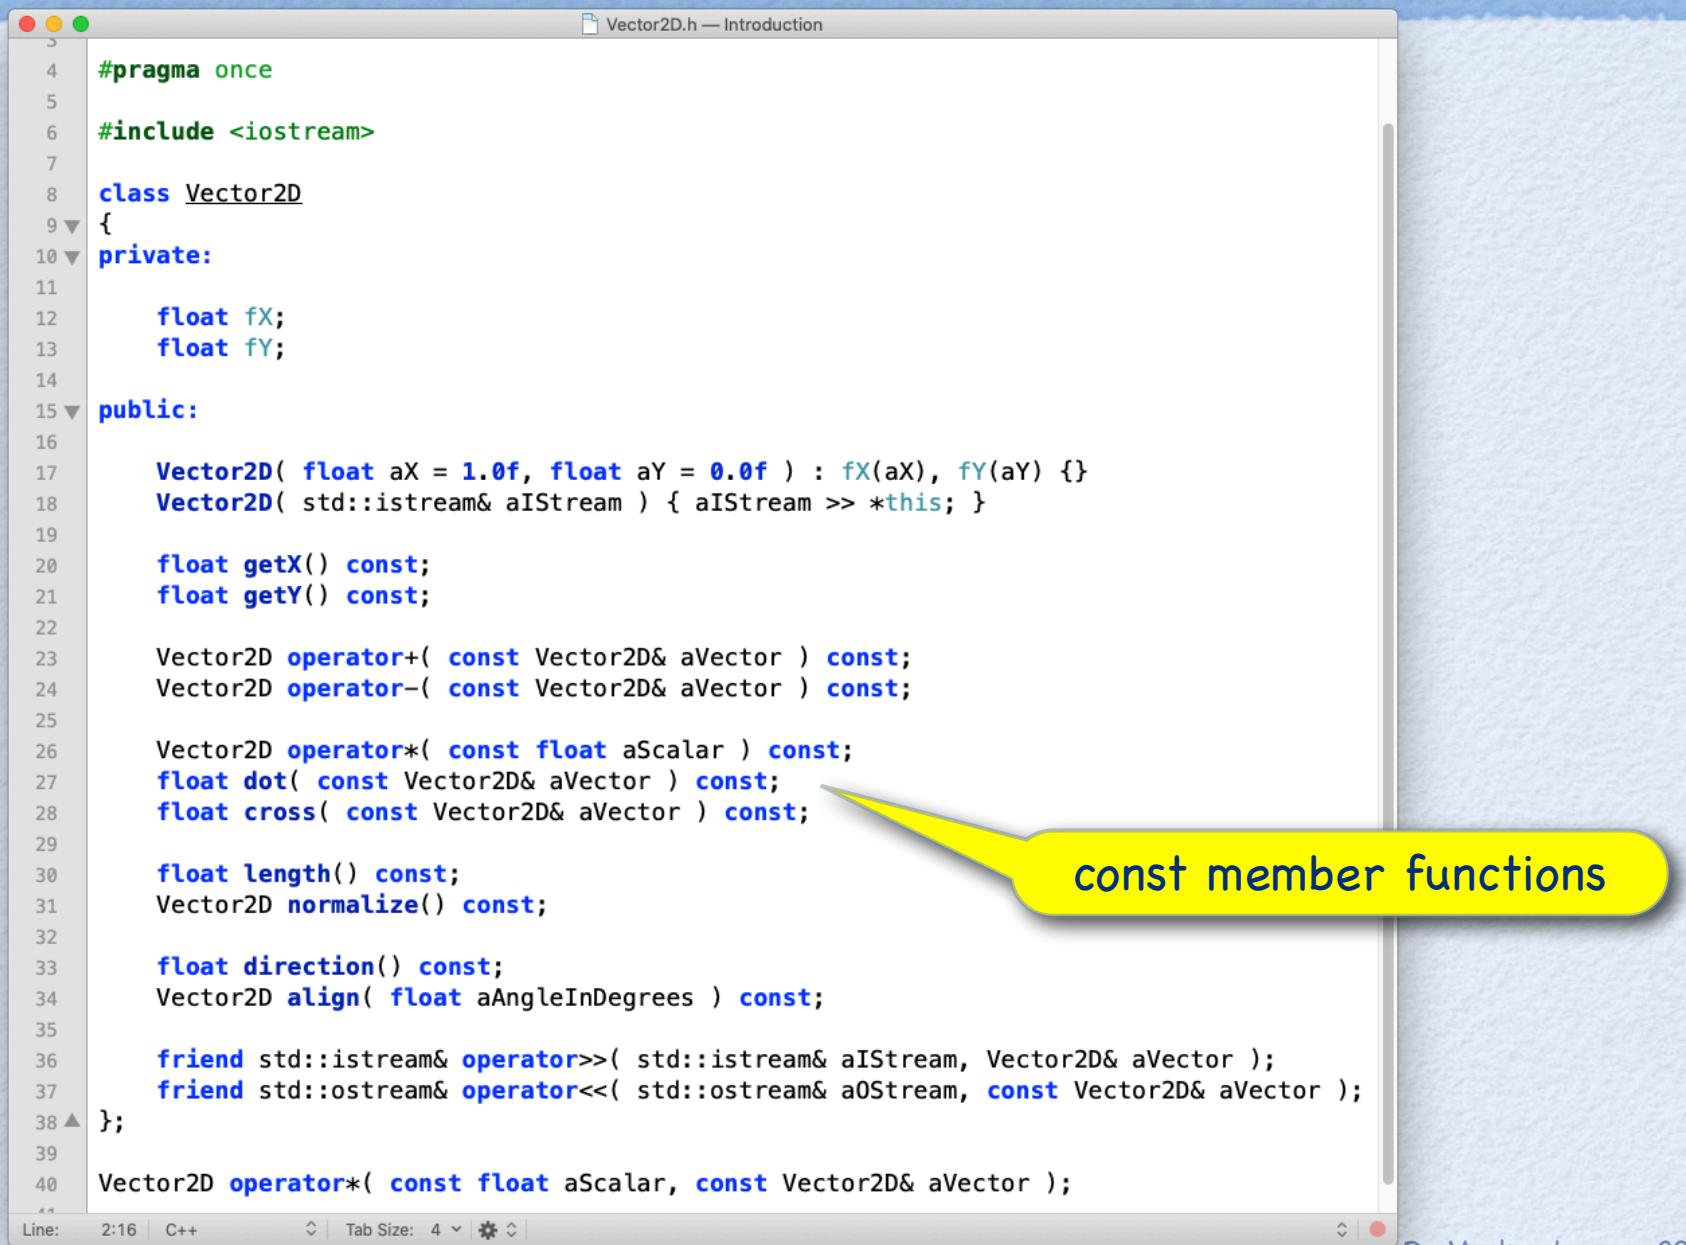
- Unlike macro definitions, const objects have an address!

```
#define BUF_SIZE 128 // macro definition
```

```
const int BufferSize = BUF_SIZE; // initialized at compile time
```

Everything that should or
must not change is marked
with the **const** keyword.

Initialized Vector2D objects are read-only.



```
#pragma once
#include <iostream>
class Vector2D
{
private:
    float fX;
    float fY;
public:
    Vector2D( float aX = 1.0f, float aY = 0.0f ) : fX(aX), fY(aY) {}
    Vector2D( std::istream& aInputStream ) { aInputStream >> *this; }

    float getX() const;
    float getY() const;

    Vector2D operator+( const Vector2D& aVector ) const;
    Vector2D operator-( const Vector2D& aVector ) const;

    Vector2D operator*( const float aScalar ) const;
    float dot( const Vector2D& aVector ) const;
    float cross( const Vector2D& aVector ) const;

    float length() const;
    Vector2D normalize() const;

    float direction() const;
    Vector2D align( float aAngleInDegrees ) const;

    friend std::istream& operator>>( std::istream& aInputStream, Vector2D& aVector );
    friend std::ostream& operator<<( std::ostream& aOutputStream, const Vector2D& aVector );
};

Vector2D operator*( const float aScalar, const Vector2D& aVector );
```

const member functions

References (C++-98)

- A reference introduces a new name (alias) for an existing object:

```
int BlockSize = 512;
```

```
int& BufferSize = BlockSize;
```



BufferSize is an
alias for BlockSize

Constant References (C++-98)

- A constant reference yields a new name for a constant object:

```
const int FixedBlockSize = 512;
```

```
const int& FixedBufferSize = FixedBlockSize;
```

- A constant reference defines an immutable alias to an existing object.

References prevent copies from being made.

Reference Parameters (C++-98)

- C++ uses call-by-value as default parameter passing mechanism.

```
void Assign( int aPar, int aVal ) { aPar = aVal; }

Assign( val, 3 );           // val unchanged
```

- A reference parameter yields call-by-reference:

```
void AssignR( int& aPar, int aVal ) { aPar = aVal; }

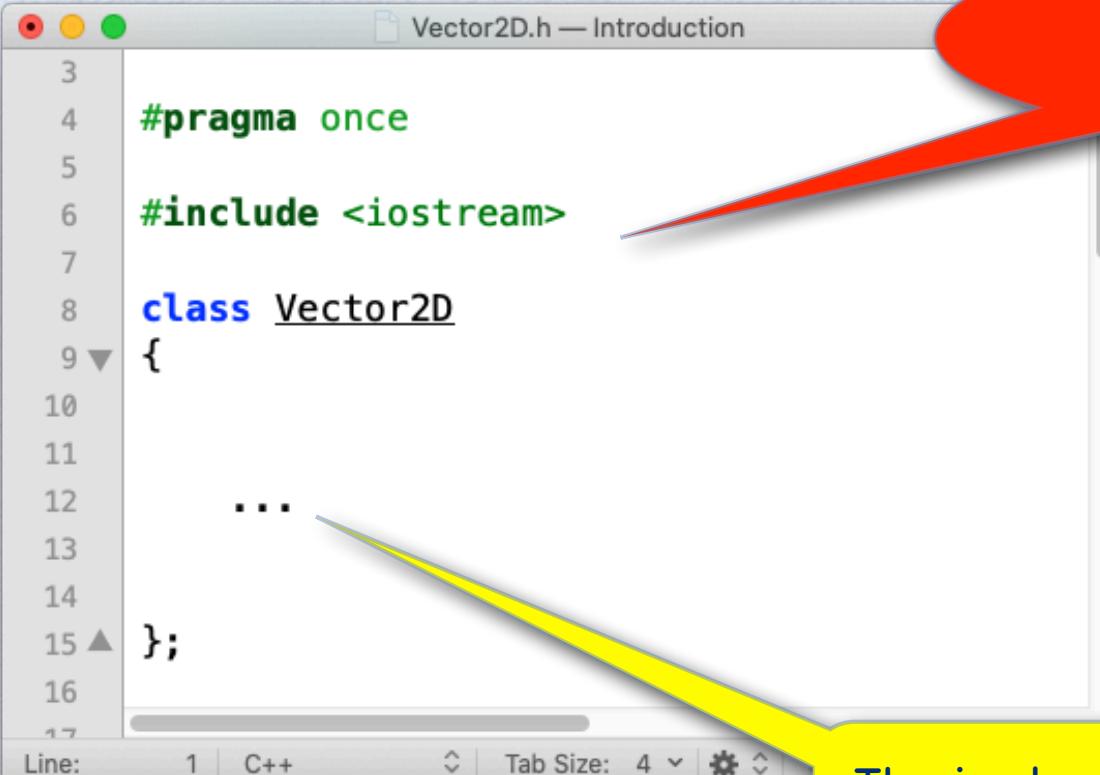
AssignR( val, 3 );          // val is set to 3
```

- A const reference parameter yields call-by-reference, but the value of the parameter is read-only:

```
void AssignCR( const int& aPar, int aVal ) { aPar = aVal; } // error
```

Class Implementation

Include File: Vector2D.h



```
3
4 #pragma once
5
6 #include <iostream>
7
8 class Vector2D
9 {
10
11
12     ...
13
14 };
15
16
17
```

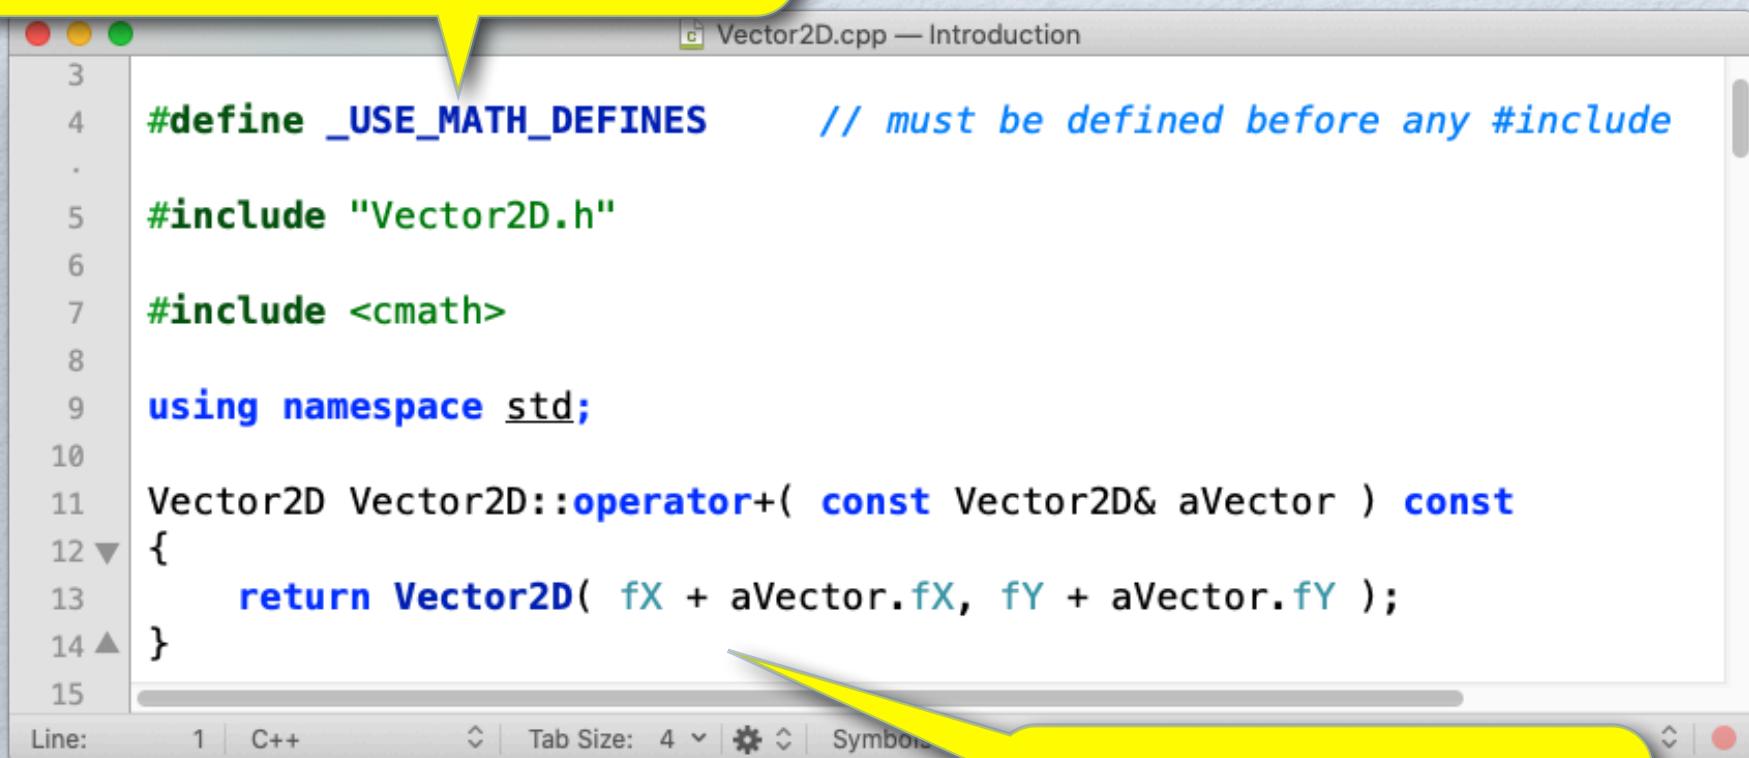
The code editor window shows the beginning of a C++ class definition. The file is titled "Vector2D.h — Introduction". The code includes a `#pragma once` directive, an `#include <iostream>` directive, and a class definition starting with `class Vector2D`. The implementation part of the class, indicated by the ellipsis (`...`), is shown in a yellow callout bubble.

We do not select any namespace yet!

The implementation goes to Vector2D.cpp

Implementation File: Vector2D.cpp

Macro definition to make math definitions available (e.g., M_PI)



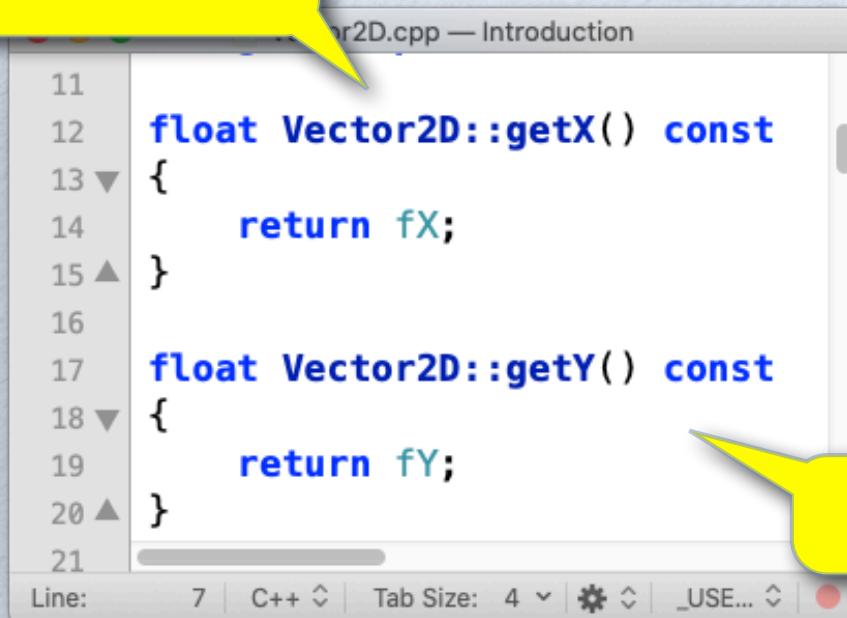
```
3
4 #define _USE_MATH_DEFINES      // must be defined before any #include
5
6
7 #include "Vector2D.h"
8
9 #include <cmath>
10
11 using namespace std;
12
13 Vector2D Vector2D::operator+( const Vector2D& aVector ) const
14 {
15     return Vector2D( fX + aVector.fX, fY + aVector.fY );
}
```

Implementation

Member Implementation

- When implementing a member function of a class in C++ we must explicitly specify the class name using a scope identifier within the signature of the member function.
- A scope identifier is a name followed by two colons (e.g. `Vector2D::`).

Scope identifier `Vector2D::`



A screenshot of a code editor window titled "Vector2D.cpp — Introduction". The code editor displays two member functions of the `Vector2D` class:

```
11
12 float Vector2D::getX() const
13 {
14     return fx;
15 }
16
17 float Vector2D::getY() const
18 {
19     return fy;
20 }
```

The scope identifier `Vector2D::` is highlighted in blue. A yellow callout bubble points to the identifier with the text "Scope identifier `Vector2D::`". Another yellow callout bubble points to the `getX()` and `getY()` functions with the text "Getters for Vector2D".

C++ Code Organization

- Classes are **defined** in include files (i.e., .h).
- Class members are **implemented** in source files (i.e., .cpp).
- There are **exceptions** when working with templates.

Standard Boilerplate Code

Guard
against
repeated
inclusion

```
#ifndef HEADER_H_
#define HEADER_H_

/* Body of Header */

#endif /* HEADER_H_ */
```

#pragma once (Visual Studio)

```
#pragma once
```

Guard against
repeated inclusion

```
/* Body of Header */
```

Object Initialization

- A **class** defines a new **data type**. Instances of this data type are **objects** that need initialization.
- Each class defines explicitly (or implicitly) some special member functions, called **constructors**, that are executed whenever we create new objects of a class.
- The job of a constructor is to ensure that the data members of each objects are set to some sensible initial values.

Constructors

- Constructors may be overloaded.
- The concrete constructor arguments determine which constructor to use.
- Constructors are executed automatically whenever a new object is created.

Default Constructor

- A default constructor is one that does not take any arguments.
- The compiler will synthesize a default constructor, when no other constructors have been specified.
- There are situations when the compiler needs to create objects in an environment agnostic way (e.g., array of objects). In this case the compiler relies on the existence of the default constructor.
- If some data members have built-in or compound types, then the class should not rely on the synthesized default constructor!

Class Vector2D - Constructors

```
#pragma once
#include <iostream>

class Vector2D
{
private:
    float fX;
    float fY;

public:
    Vector2D( float aX = 1.0f, float aY = 0.0f ) : fX(aX), fY(aY) {}
    Vector2D( std::istream& aInputStream ) { aInputStream >> *this; }

    float getX() const;
    float getY() const;

    Vector2D operator+( const Vector2D& aVector ) const;
    Vector2D operator-( const Vector2D& aVector ) const;

    Vector2D operator*( const float aScalar ) const;
    float dot( const Vector2D& aVector ) const;
    float cross( const Vector2D& aVector ) const;

    float length() const;
    Vector2D normalize() const;

    float direction() const;
    Vector2D align( float aAngleInDegrees ) const;

    friend std::istream& operator>>( std::istream& aInputStream, Vector2D& aVector );
    friend std::ostream& operator<<( std::ostream& aOutputStream, const Vector2D& aVector );
};

Vector2D operator*( const float aScalar, const Vector2D& aVector );
```

Line: 2:16 | C++ | Tab Size: 4 |

Default argument

Note, we use inlined trivial constructors here.

Constructor Initializer

- A constructor initializer is a comma-separated list of member initializers, which is declared between the signature of the constructor and its body.
- Constructor initializers take the form of function calls where the name of the function coincides with name of the instance variable being initialized.

```
Vector2D( float aX = 1.0f, float aY = 0.0f ) : fX(aX), fY(aY) {}
```

Implicit Class-Type Conversions

- A constructor that can be called with a single argument defines an implicit conversion from the parameter type to the class type.

```
Vector2D( std::istream& aIStream ) { aIStream >> *this; }
```

Conversion from std::istream to Vector2D

- Accordingly, we can use an `istream` where an object of type `Vector2D` is expected:

```
Vector2D( std::istream& aIStream ) { aIStream >> *this; }
```

```
Vector2D operator+( const Vector2D& aVector ) const;
```

```
Vector2D result = vec + cin;
```

`std::istream` → `Vector2D`

Suppressing Implicit Conversions

- When a constructor is declared **explicit**, the compiler will **not** use it as a conversion operator.

```
explicit Vector2D( std::istream& aIStream ) { aIStream >> *this; }
```

- Accordingly, we **cannot** use an **istream** where an object of type **Vector2D** is expected:

```
explicit Vector2D( std::istream& aIStream ) { aIStream >> *this; }
```

```
Vector2D operator+( const Vector2D& aVector ) const;
```



```
Vector2D result = vec + cin;
```

Error

**It is application-dependent
whether implicit conversion
needs to be supported.**

Rule of Thumb

- Keep C++ headers free of implementations, unless you want the implementations to be inlined or they are templates.

Basic 2D vector operations

Operator Overloading

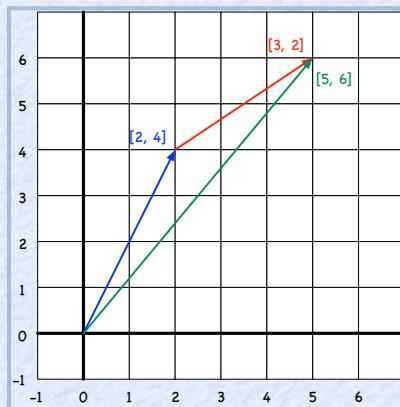
- C++ supports operator overloading.
- Overloaded operators are like normal functions, but are defined using a pre-defined operator symbol.
- You cannot change the priority and associativity of an operator.
- Operators are selected by the compiler based on the static types of the specified operands.

Member Operators vs Ad hoc Operators

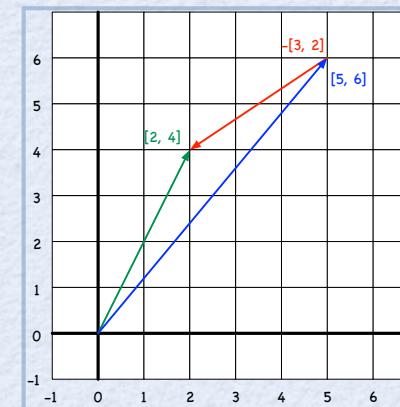
- There are two forms of operator overloading:
 - member operator
 - ad hoc operator
- A member operator receives a first implicit argument - this object. In other words, an overloaded operator in C++ is like any other non-static method function defined for a class. Non-static member functions receive as first argument "this" object. As a consequence, we only specify the second argument for binary operators like '+' or '-'.
- Ad hoc operators are not members of a class. Their signature has to match the signature of the operator to be defined. That is, an ad hoc definition of a binary operator '+' requires two arguments: the left-hand side of '+' and the right-hand side of '+'.

Addition & Subtraction

$$[5,6] = [2,4] + [3,2]$$



$$[2,4] = [5,6] - [3,2]$$



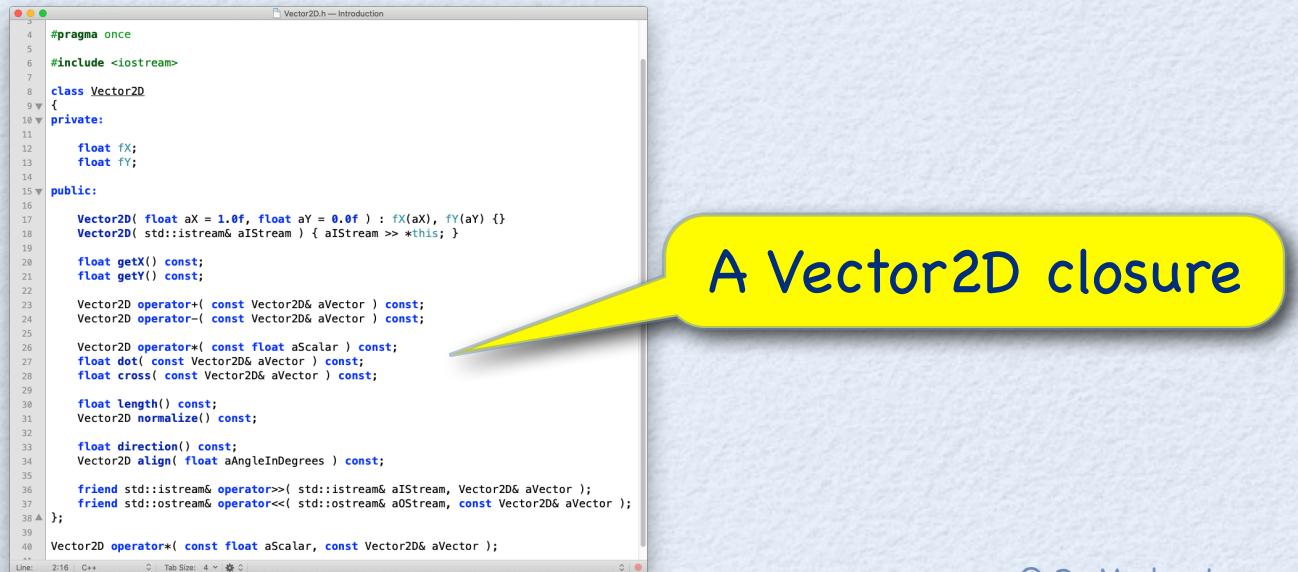
```
Vector2D Vector2D::operator+( const Vector2D& aVector ) const
{
    return Vector2D( fX + aVector.fX, fY + aVector.fY );
}

Vector2D Vector2D::operator-( const Vector2D& aVector ) const
{
    return Vector2D( fX - aVector.fX, fY - aVector.fY );
}
```

Member operators

Closures – Lexical Scoped Name Binding

- Classes provide a lexical scoped name binding for methods. This allows methods of a class to access instance variables and other methods of the class that are defined outside the scope of a given method.
- Operationally, classes create closures that are records storing a function (or a set of functions) together with an environment. The environment associates each free variable of the function(s) to values to which the name was bound when the closure was created.



```
#pragma once
#include <iostream>
class Vector2D
{
private:
    float fx;
    float fy;
public:
    Vector2D( float aX = 1.0f, float aY = 0.0f ) : fx(aX), fy(aY) {}
    Vector2D( std::istream& aInputStream ) { aInputStream >> *this; }

    float getX() const;
    float getY() const;

    Vector2D operator+( const Vector2D& aVector ) const;
    Vector2D operator-( const Vector2D& aVector ) const;

    Vector2D operator*( const float aScalar ) const;
    float dot( const Vector2D& aVector ) const;
    float cross( const Vector2D& aVector ) const;

    float length() const;
    Vector2D normalize() const;

    float direction() const;
    Vector2D align( float aAngleInDegrees ) const;

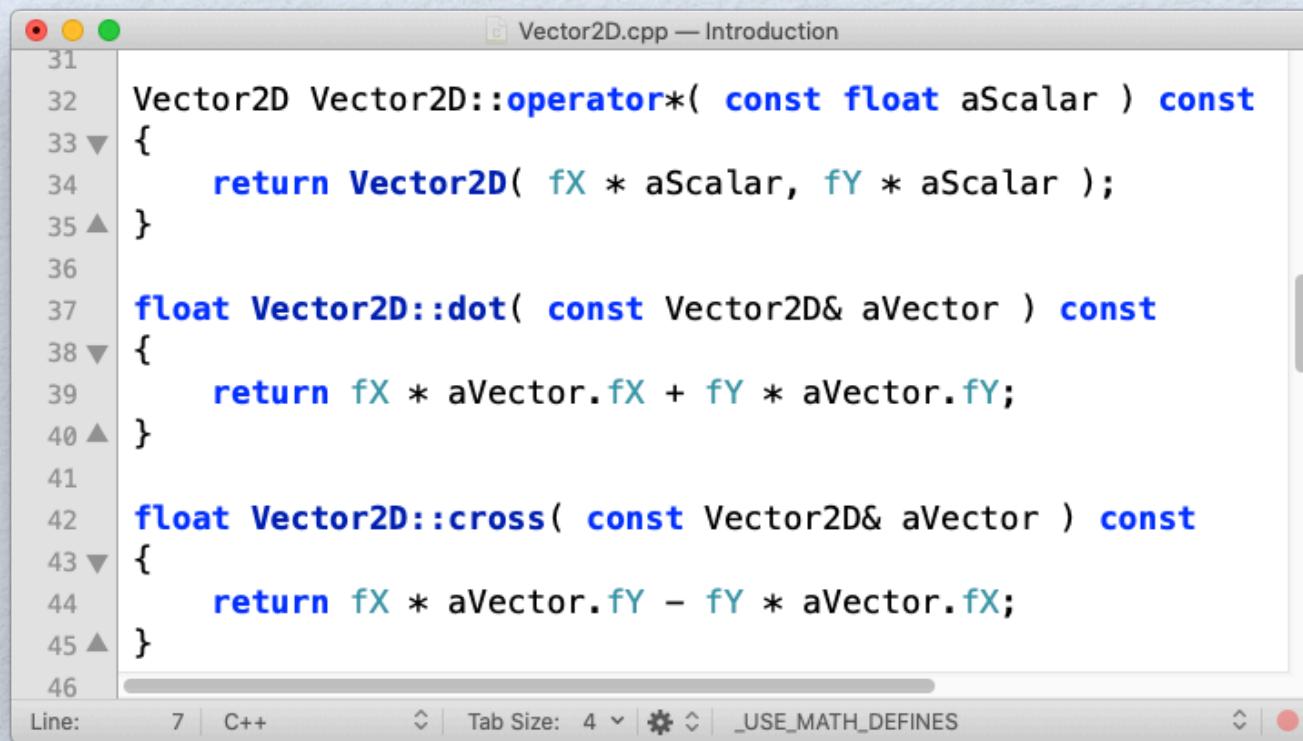
    friend std::istream& operator>>( std::istream& aInputStream, Vector2D& aVector );
    friend std::ostream& operator<<( std::ostream& aOutputStream, const Vector2D& aVector );
};

Vector2D operator*( const float aScalar, const Vector2D& aVector );
```

A Vector2D closure

Scalar Multiplication, Dot Product, and Cross Product

- We use scalar multiplication to scale a vector uniformly.
- The dot product (inner product) is a measure of the difference between the directions in which the two vectors point.
- The 2D cross product yields a scalar that we use it to determine whether consecutive line segments turn left or right.



The screenshot shows a code editor window titled "Vector2D.cpp — Introduction". The code defines three member functions for the `Vector2D` class:

```
31
32     Vector2D Vector2D::operator*( const float aScalar ) const
33 {
34     return Vector2D( fX * aScalar, fY * aScalar );
35 }
36
37     float Vector2D::dot( const Vector2D& aVector ) const
38 {
39     return fX * aVector.fX + fY * aVector.fY;
40 }
41
42     float Vector2D::cross( const Vector2D& aVector ) const
43 {
44     return fX * aVector.fY - fY * aVector.fX;
45 }
```

The code uses standard C++ syntax with `const` qualifiers and floating-point arithmetic. The `operator*` overload returns a new `Vector2D` object. The `dot` method calculates the dot product of the current vector with another. The `cross` method calculates the 2D cross product with another vector, which is equivalent to the determinant of a 2x2 matrix formed by the two vectors.

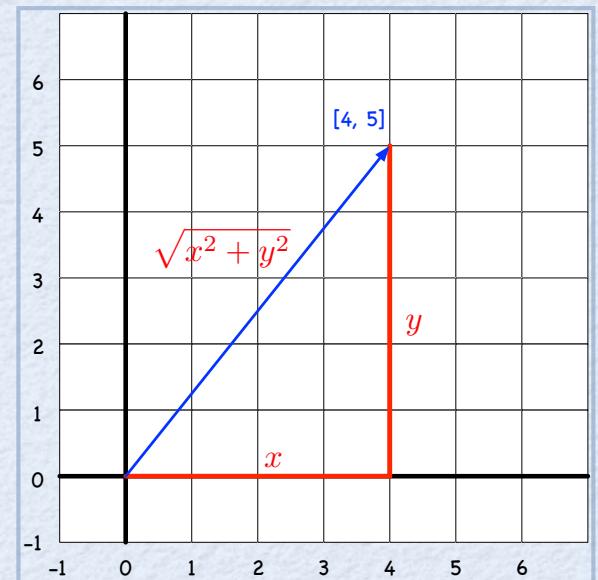
Vector Length and Unit Vector of a Vector

- The **length** of a vector (magnitude) is the hypotenuse of the right-angled triangle formed by the vector coordinates x and y.
- The **unit vector** of a vector is a vector with length 1. (In the code below `*this` refers to the `this` object, that is, the vector object for which we calculate the unit vector.)

```
Vector2D.cpp — Introduction
47 float Vector2D::length() const
48 {
49     float val = sqrt(fX * fX + fY * fY);
50
51     return round( val * 100.0f ) / 100.0f;
52 }
53
54 Vector2D Vector2D::normalize() const
55 {
56     return *this * (1.0f/length());
57 }
```

The screenshot shows a code editor window titled "Vector2D.cpp — Introduction". The code defines two methods for the "Vector2D" class. The first method, "length()", calculates the magnitude of the vector using the Pythagorean theorem ($\sqrt{fX^2 + fY^2}$) and rounds the result to two decimal places. The second method, "normalize()", returns a new vector with the same direction but a length of 1.0, achieved by dividing the current vector by its length.

vector length



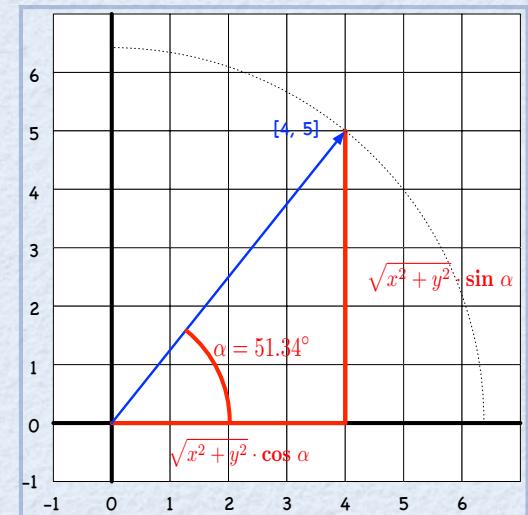
Direction and Align

- The direction of a vector is the arctangent of the right-angled triangle formed by the vector coordinates x and y .
- We can align/rotate a vector, without changing its length, by multiplying its length with the sine and the cosine of the direction angle to obtain the new x and y coordinates, respectively.

C++ cast M_PI to float

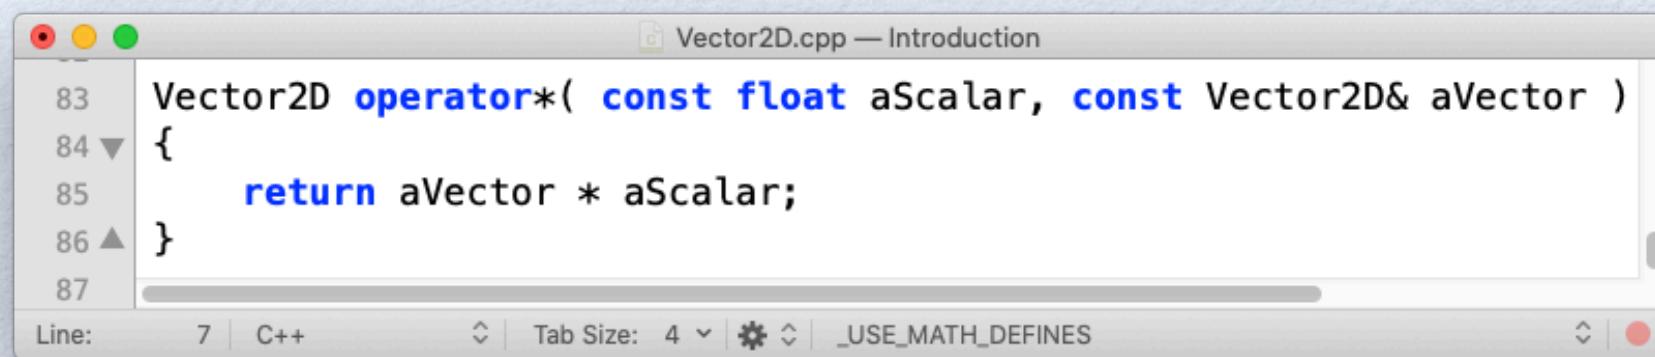
```
Line: 11 | C++ | Tab Size: 4 | std
58
59     float Vector2D::direction() const
60 {
61     float val = atan2( fY, fX ) * 180.0f / static_cast<float>(M_PI);
62
63     return round( val * 100.0f ) / 100.0f;
64 }
65
66 Vector2D Vector2D::align( float aAngleInDegrees ) const
67 {
68     float lRadians = aAngleInDegrees * static_cast<float>(M_PI) / 180.0f;
69
70     return length() * Vector2D( cos( lRadians ), sin( lRadians ) );
71 }
72
```

direction/align



Ad hoc Operator *

- The member operator for scalar multiplication only allows for `vector * scalar`. However, multiplication is commutative, that is changing the order of the operands does not change the result.
- We can recover the commutativity of scalar multiplication by defining an ad hoc multiplication operator that takes a scalar as first argument and a vector as the second:



The screenshot shows a code editor window titled "Vector2D.cpp — Introduction". The code is as follows:

```
83 Vector2D operator*( const float aScalar, const Vector2D& aVector )
84 {
85     return aVector * aScalar;
86 }
87
```

The code defines a member function `operator*` for the `Vector2D` class. It takes a `const float` parameter `aScalar` and a `const Vector2D&` parameter `aVector`. The function returns the result of multiplying `aVector` by `aScalar`. The code editor interface includes line numbers, syntax highlighting, and standard toolbars at the bottom.

Data I/O

I/O in C++ is operator based.



```
SimpleIO.cpp
2 #include <iostream>
3
4 using namespace std;
5
6 int main()
7 {
8     cout << "Enter two numbers:" << endl;
9     int v1, v2;
10    cin >> v1 >> v2;
11    cout << "The sum of " << v1 << " and " << v2
12    << " is " << v1 + v2 << endl;
13
14    return 0;
15 }
16
```

Line: 17 Column: 1 C++ Tab Size: 4 main

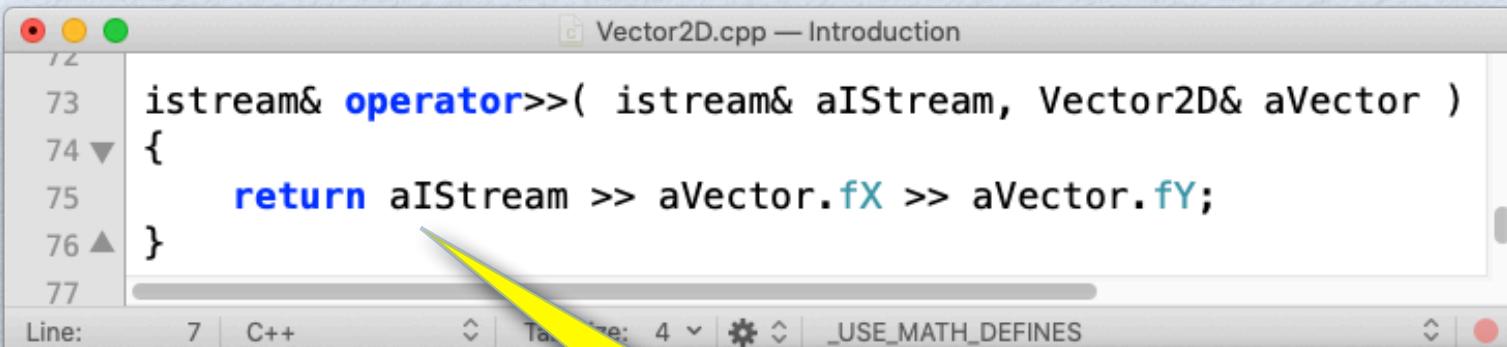


```
COS30008
Kamala:COS30008 Markus$ g++ -o SimpleIO SimpleIO.cpp
Kamala:COS30008 Markus$ ./SimpleIO
Enter two numbers:
7
5
The sum of 7 and 5 is 12
Kamala:COS30008 Markus$
```

- C++ relies on the binary operators `<<` and `>>` to perform formatted input and output, respectively.

The Vector2D Input Operator >>

- The input operator performs formatted input of the vector coordinates, that is, we try to fetch two floating point values from the input stream.
- If the input fails, either of the coordinates or both are set to 0.0.
- At the end, the input operator returns the stream object to allow for a “chaining” of input operations.



```
Vector2D.cpp — Introduction
73 istream& operator>>( istream& aIStream, Vector2D& aVector )
74 {
75     return aIStream >> aVector.fX >> aVector.fY;
76 }
77
```

A screenshot of a Mac OS X-style IDE window titled "Vector2D.cpp — Introduction". The code editor shows a portion of a C++ file with the following content:

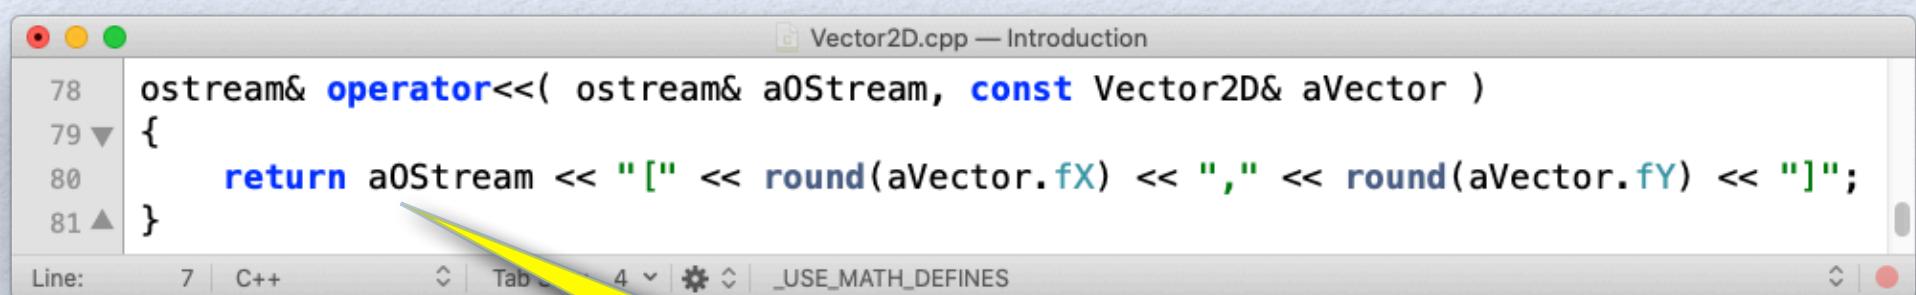
```
73 istream& operator>>( istream& aIStream, Vector2D& aVector )
74 {
75     return aIStream >> aVector.fX >> aVector.fY;
76 }
77
```

The line numbers 73 through 77 are visible on the left. A yellow callout bubble points from the word "return" in line 75 to the text "Return reference to input stream." at the bottom right.

Return reference to input stream.

The Vector2D Output Operator <<

- The output operator assigns a “textual representation” to objects of type Vector2D.
- It performs formatted output, that is, the coordinate values are rendered as integer values in the output. The round function yields the nearest integral value.
- At the end, the input operator returns the stream object to allow for a “chaining” of output operations.



```
Vector2D.cpp — Introduction
78 ostream& operator<<( ostream& aOutputStream, const Vector2D& aVector )
79 {
80     return aOutputStream << "[" << round(aVector.fX) << "," << round(aVector.fY) << "]";
81 }
```

A screenshot of a C++ code editor window titled "Vector2D.cpp — Introduction". The code shows the implementation of the output operator (<<) for the Vector2D class. The operator takes an output stream (ostream&) and a Vector2D object (const Vector2D&) as parameters. It returns the output stream itself after printing the vector's coordinates in brackets. A yellow callout points to the "return" statement with the text "Return reference to output stream."

Return reference to output stream.

Class Vector2D - The Friends

```
#pragma once
#include <iostream>
class Vector2D
{
private:
    float fX;
    float fY;
public:
    Vector2D( float aX = 1.0f, float aY = 0.0f ) : fX(aX), fY(aY) {}
    Vector2D( std::istream& aInputStream ) { aInputStream >> *this; }

    float getX() const;
    float getY() const;

    Vector2D operator+( const Vector2D& aVector ) const;
    Vector2D operator-( const Vector2D& aVector ) const;

    Vector2D operator*( const float aScalar ) const;
    float dot( const Vector2D& aVector ) const;
    float cross( const Vector2D& aVector ) const;

    float length() const;
    Vector2D normalize() const;

    float direction() const;
    Vector2D align( float aAngleInDegrees ) const;

    friend std::istream& operator>>( std::istream& aInputStream, Vector2D& aVector );
    friend std::ostream& operator<<( std::ostream& aOutputStream, const Vector2D& aVector );
};

Vector2D operator*( const float aScalar, const Vector2D& aVector );
```

Friends

- Friends are allowed to access private members of classes.
- A class declares its friends explicitly.
- Friends enable uncontrolled access to members.
- The friend mechanism induces a particular programming (C++) style.
- The friend mechanism is not object-oriented!
- I/O depends on the friend mechanism.

The Friend Mechanism

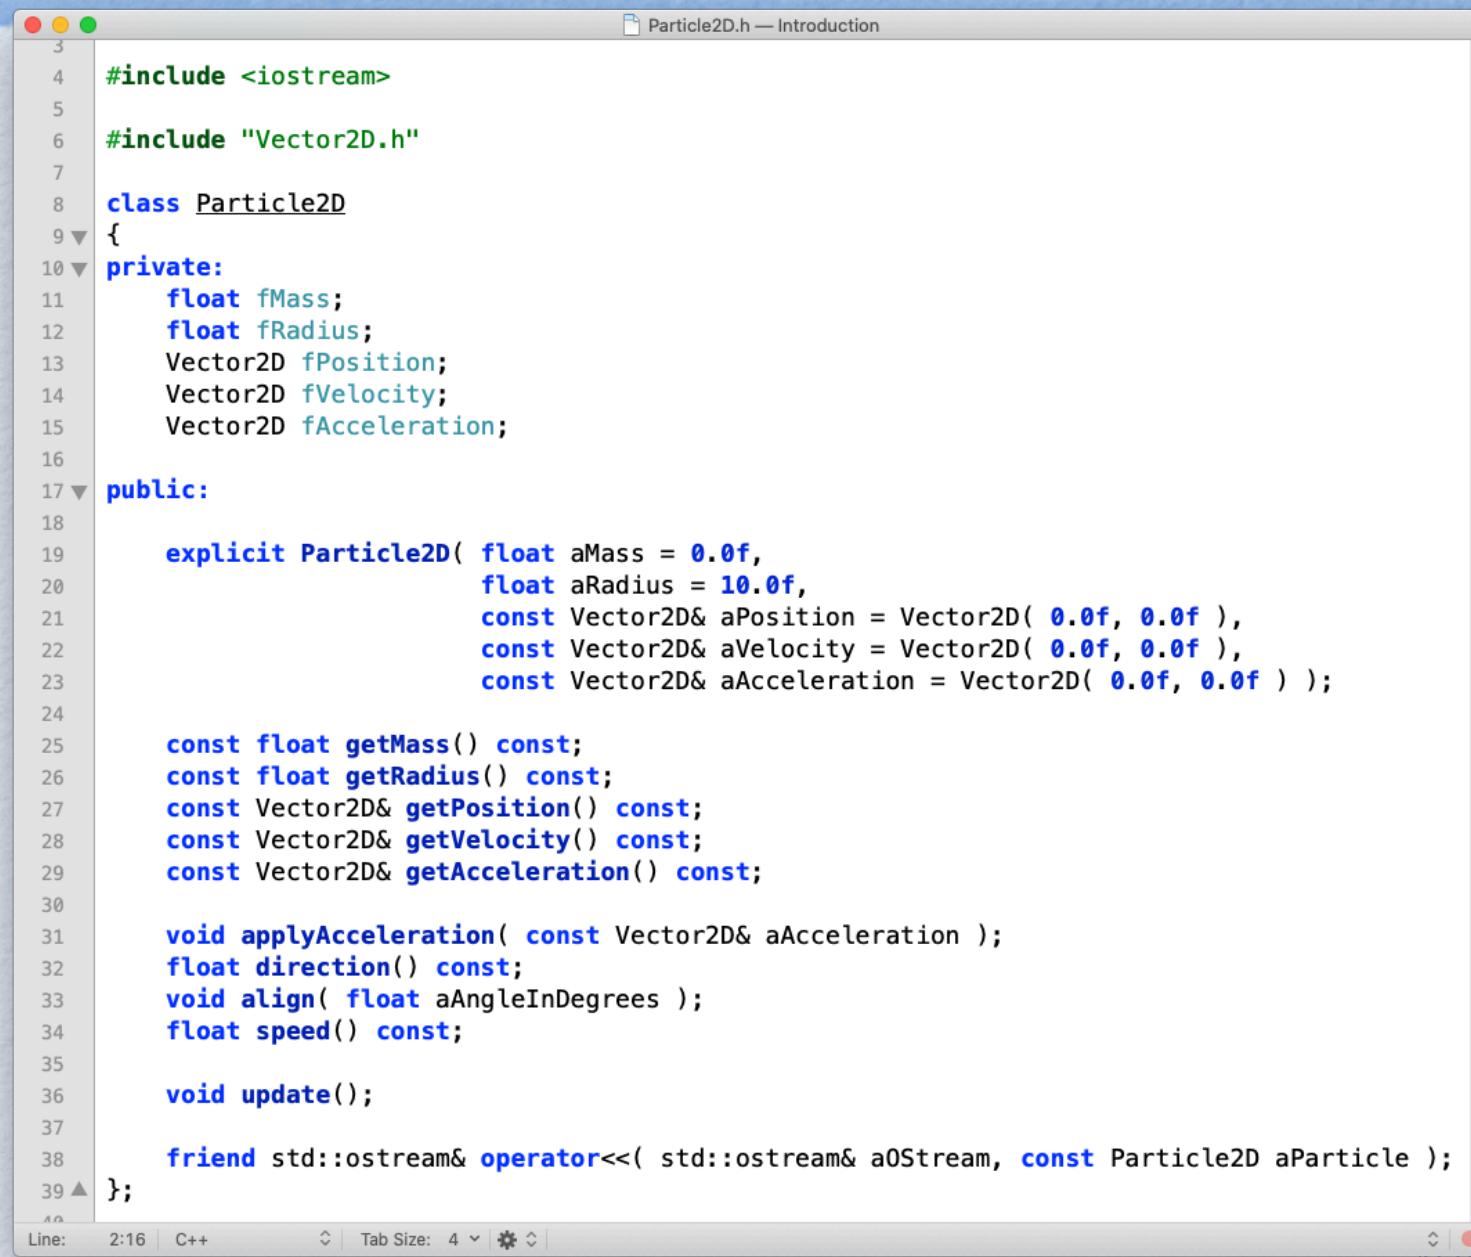
Friends are self-contained procedures (or functions) that do not belong to a specific class, but have access to the members of a class, when this class declares those procedures as friends.

Class Composition: Particle2D

Requirements for a particle class

- Particles represent physical model entities with the following attributes:
 - Mass (the particle's resistance – inertia – to change in its velocity)
 - Radius (particles can be rendered as filled circles)
 - Position, a 2D vector to assign a location in 2D space.
 - Velocity, a 2D vector to capture the speed of a particle in a given direction
 - Acceleration, a 2D vector representing the acceleration of a particle in a given direction due to the application of some force (e.g., gravity)
- References
 - Leonard Susskind and George Hrabovsky: The Theoretical Minimum – What You Need to Know to Start Doing Physics. Allen Lane (2013)
 - www.codingmath.com

Particle2D: A Simple 2D Particle Class



The screenshot shows a Mac OS X application window titled "Particle2D — Introduction". The window contains the source code for the Particle2D class. The code is color-coded: comments are in green, keywords are in blue, and variable names are in black. The code defines a class with private members (mass, radius, position, velocity, acceleration) and public methods (constructor, getters for mass and radius, setters for position, velocity, and acceleration, applyAcceleration, direction, align, speed, update, and a friend operator for output).

```
3 #include <iostream>
4
5 #include "Vector2D.h"
6
7 class Particle2D
8 {
9 private:
10     float fMass;
11     float fRadius;
12     Vector2D fPosition;
13     Vector2D fVelocity;
14     Vector2D fAcceleration;
15
16 public:
17     explicit Particle2D( float aMass = 0.0f,
18                           float aRadius = 10.0f,
19                           const Vector2D& aPosition = Vector2D( 0.0f, 0.0f ),
20                           const Vector2D& aVelocity = Vector2D( 0.0f, 0.0f ),
21                           const Vector2D& aAcceleration = Vector2D( 0.0f, 0.0f ) );
22
23     const float getMass() const;
24     const float getRadius() const;
25     const Vector2D& getPosition() const;
26     const Vector2D& getVelocity() const;
27     const Vector2D& getAcceleration() const;
28
29     void applyAcceleration( const Vector2D& aAcceleration );
30     float direction() const;
31     void align( float aAngleInDegrees );
32     float speed() const;
33
34     void update();
35
36     friend std::ostream& operator<<( std::ostream& aOutputStream, const Particle2D aParticle );
37
38 };
39
```

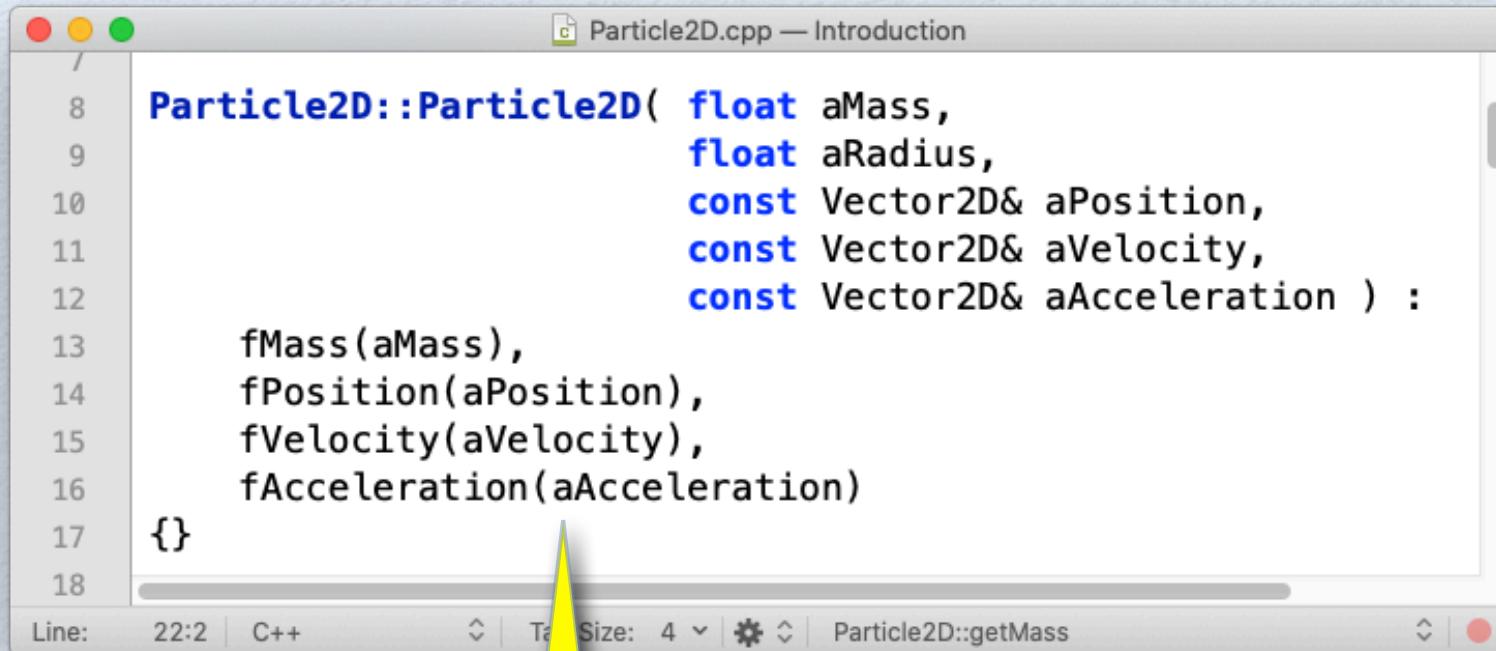
Line: 2:16 | C++ | Tab Size: 4 |  | 84 | © Dr. Markus Lumpe, 2024

Explicit Particle2D Constructor

```
3 #include <iostream>
4
5 #include "Vector2D.h"
6
7 class Particle2D
8 {
9 private:
10     float fMass;
11     float fRadius;
12     Vector2D fPosition;
13     Vector2D fVelocity;
14     Vector2D fAcceleration;
15
16 public:
17     explicit Particle2D( float aMass = 0.0f,
18                           float aRadius = 10.0f,
19                           const Vector2D& aPosition = Vector2D( 0.0f, 0.0f ),
20                           const Vector2D& aVelocity = Vector2D( 0.0f, 0.0f ),
21                           const Vector2D& aAcceleration = Vector2D( 0.0f, 0.0f ) );
22
23     const float getMass() const;
24     const float getRadius() const;
25     const Vector2D& getPosition() const;
26     const Vector2D& getVelocity() const;
27     const Vector2D& getAcceleration() const;
28
29     void applyAcceleration( const Vector2D& aAcceleration );
30     float direction() const;
31     void align( float aAngleInDegrees );
32     float speed() const;
33
34     void update();
35
36     friend std::ostream& operator<<( std::ostream& aOutputStream, const Particle2D aParticle );
37
38 };
39
```

The constructor is marked `explicit` to prevent automatic type conversion from `float` to `Particle2D` when all other parameters are omitted (use default).

Constructor Implementation



```
Particle2D::Particle2D( float aMass,
                        float aRadius,
                        const Vector2D& aPosition,
                        const Vector2D& aVelocity,
                        const Vector2D& aAcceleration ) :
    fMass(aMass),
    fPosition(aPosition),
    fVelocity(aVelocity),
    fAcceleration(aAcceleration)
{}
```

We use a constructor initializer list to prevent a default initialization of the Vector2D member variables. Instead, we employ the compiler-generated copy constructor to initialize the Vector2D member variables. (The Qt framework uses this approach extensively.)

The Particle2D Operations

A screenshot of a C++ code editor showing the file `Particle2D.cpp`. The code defines a class `Particle2D` with several member functions: `applyAcceleration`, `direction`, `align`, `speed`, and `update`. The `update` function is annotated with two yellow callouts: one pointing to the line `fVelocity = fVelocity + fAcceleration;` with the text "Adjust the direction of the velocity", and another pointing to the line `fPosition = fPosition + fVelocity;` with the text "Update velocity and position".

```
43
44 void Particle2D::applyAcceleration( const Vector2D& aAcceleration )
45 {
46     fAcceleration = fAcceleration + aAcceleration;
47 }
48
49 float Particle2D::direction() const
50 {
51     return fVelocity.direction();
52 }
53
54 void Particle2D::align( float aAngleInDegrees )
55 {
56     fVelocity = fVelocity.align( aAngleInDegrees );
57 }
58
59 float Particle2D::speed() const
60 {
61     return fVelocity.length();
62 }
63
64 void Particle2D::update()
65 {
66     fVelocity = fVelocity + fAcceleration;
67     fPosition = fPosition + fVelocity;
68 }
```

Line: 22:2 | C++ | Tab Size: 4 | Particle2D::getMass

The Simulation: Main

The image shows a comparison between a code editor and a terminal window.

Code Editor (Left): A screenshot of a Mac OS X-style application window titled "Main.cpp — Introduction". It contains the following C++ code:

```
#include <iostream>
#include "Particle2D.h"

using namespace std;

int main()
{
    cout << "A simple particle simulation\n" << endl;

    Particle2D obj( 0.0f,
                    10.0f,
                    Vector2D( 10.0f, 20.0f ),
                    Vector2D( 4.0f, 15.0f ),
                    Vector2D( 0.0f, -0.1f )
    );

    do
    {
        cout << obj << endl;
        obj.update();
    } while ( obj.getPosition().getY() >= 20.0f );

    cout << obj << endl;

    return 0;
}
```

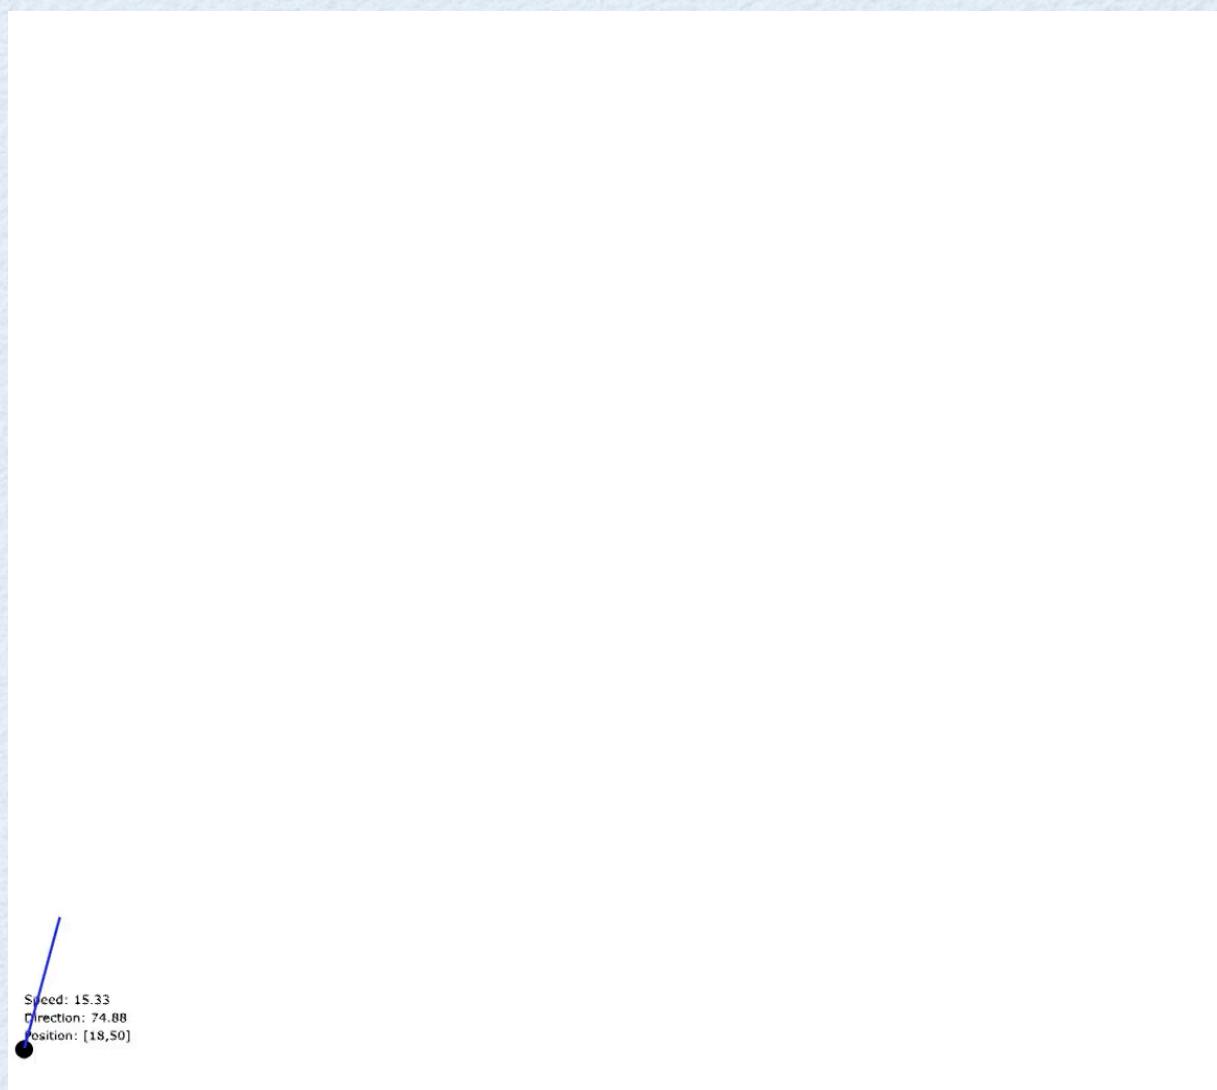
Terminal (Right): A screenshot of a Microsoft Visual Studio Debug C++ terminal window. It displays the output of the program, which is a series of particle states:

Speed	Direction	Position
14.95	-74.48	[1186,93]
15.04	-74.58	[1190,79]
15.14	-74.68	[1194,64]
15.23	-74.78	[1198,50]
15.33	-74.88	[1202,35]
15.43	-74.97	[1206,20]

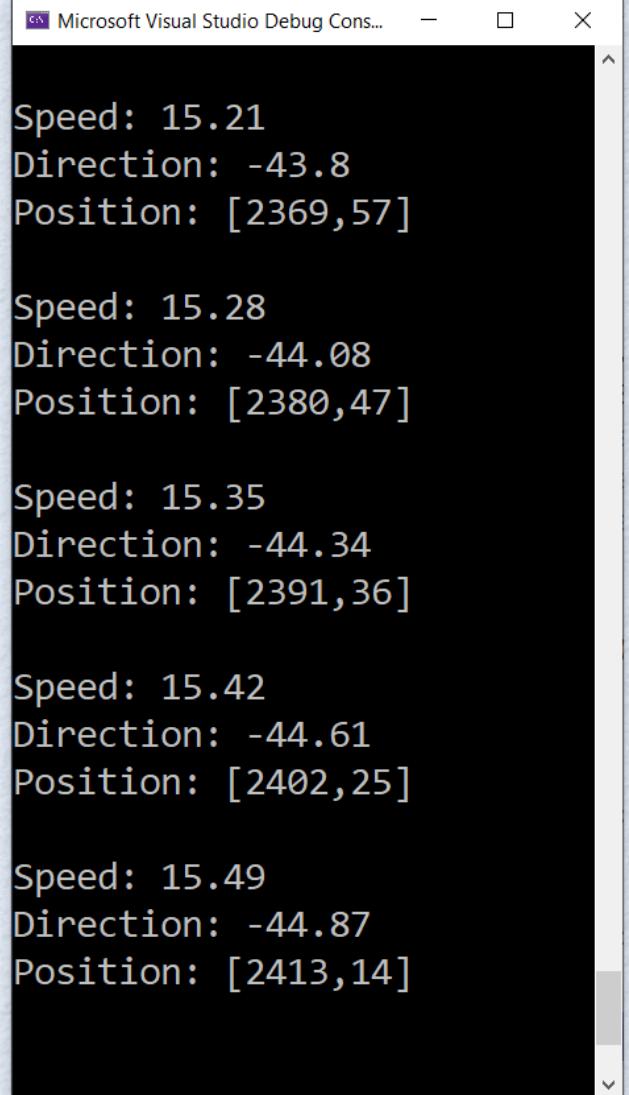
A yellow callout bubble points from the terminal window to the explanatory text below.

Loop until particle is less than 20 units above the base line.

Simulation Visualization



Best Angle: 45 Degrees



```
Line: 1 | C++ | Tab Size: 4 | ⚙ | 90
```

```
Main.cpp — Introduction

2 #include <iostream>
3
4 #include "Particle2D.h"
5
6 using namespace std;
7
8 int main()
9 {
10     cout << "A simple particle simulation\n" << endl;
11
12     Particle2D obj( 0.0f,
13                      10.0f,
14                      Vector2D( 10.0f, 20.0f ),
15                      Vector2D( 4.0f, 15.0f ),
16                      Vector2D( 0.0f, -0.1f )
17                  );
18
19     obj.align( 45.0f );    // best angle
20
21     do
22     {
23         cout << obj << endl;
24
25         obj.update();
26     } while ( obj.getPosition().getY() >= 20.0f );
27
28     cout << obj << endl;
29
30     return 0;
31 }
```

```
Speed: 15.21
Direction: -43.8
Position: [2369,57]

Speed: 15.28
Direction: -44.08
Position: [2380,47]

Speed: 15.35
Direction: -44.34
Position: [2391,36]

Speed: 15.42
Direction: -44.61
Position: [2402,25]

Speed: 15.49
Direction: -44.87
Position: [2413,14]
```

Object Models & Inheritance

C++ Object Models

- C++ supports:

- A value-based object model
- A reference-based object model

This can make programming in C++ difficult. Java, for example, has only one model - everything is a reference!

The Value-based Object Model

- Value-based object model:
 - Objects are stored on the stack.
 - Object are accessed through object variables.
 - An object's memory is implicitly released.

Valued-based Objects

- Value-based objects look and feel like records (or structs):

```
Card AceOfDiamond( Diamond, 14 );
```

```
Card TestCard( Diamond, 14 );
```

Structural
Equivalence

```
if ( TestCard == AceOfDiamond )
```

```
    cout << "The test card is " << TestCard.getName() << endl;
```

Member Selection

The Reference-based Object Model

- Reference-based object model:
 - Objects are stored on the heap.
 - Objects are accessed through pointer variables.
 - An object's memory must be explicitly released.

Reference-based Objects

- Reference-based objects require pointer variables and an explicit new and delete:

```
Card* AceOfDiamond = new Card( Diamond, 14 );
```

```
Card* TestCard = new Card( Diamond, 14 );
```

Pointer Comparison

```
if ( TestCard == AceOfDiamond )
```

```
cout << "The test card is " << TestCard->getName() << endl;
```

```
delete AceOfDiamond;
```

```
delete TestCard;
```

Member Dereference

Cardinal Rule

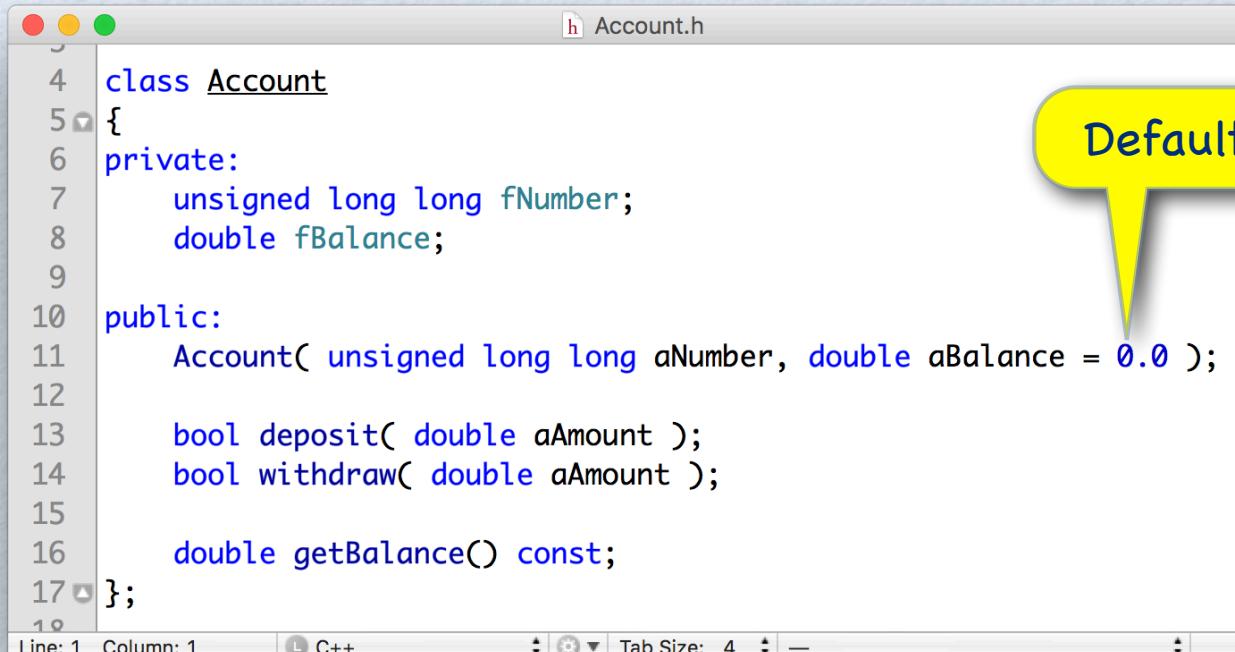
- We find both models in C++ code.
- Value semantics provides us with objects that behave like arithmetic types:

Every object is unique, but it may be known under numerous aliases (aka references).

Inheritance

- Inheritance lets us define classes that model relationships among classes, sharing what is common, and specializing only that which is inherently different.
- Inheritance is
 - A mechanism for specialization
 - A mechanism for reuse
 - Fundamental to supporting polymorphism

An Account Class



```
4 class Account
5 {
6     private:
7         unsigned long long fNumber;
8         double fBalance;
9
10    public:
11        Account( unsigned long long aNumber, double aBalance = 0.0 );
12
13        bool deposit( double aAmount );
14        bool withdraw( double aAmount );
15
16        double getBalance() const;
17    };
18
```

Line: 1 Column: 1 | L C++ | Tab Size: 4 | —

Default argument

- An account has a number ($2^{64} - 1$ values) and a balance.
- To create an account we need a number. Funds can be credited to an account once it has been created.

Account Implementation

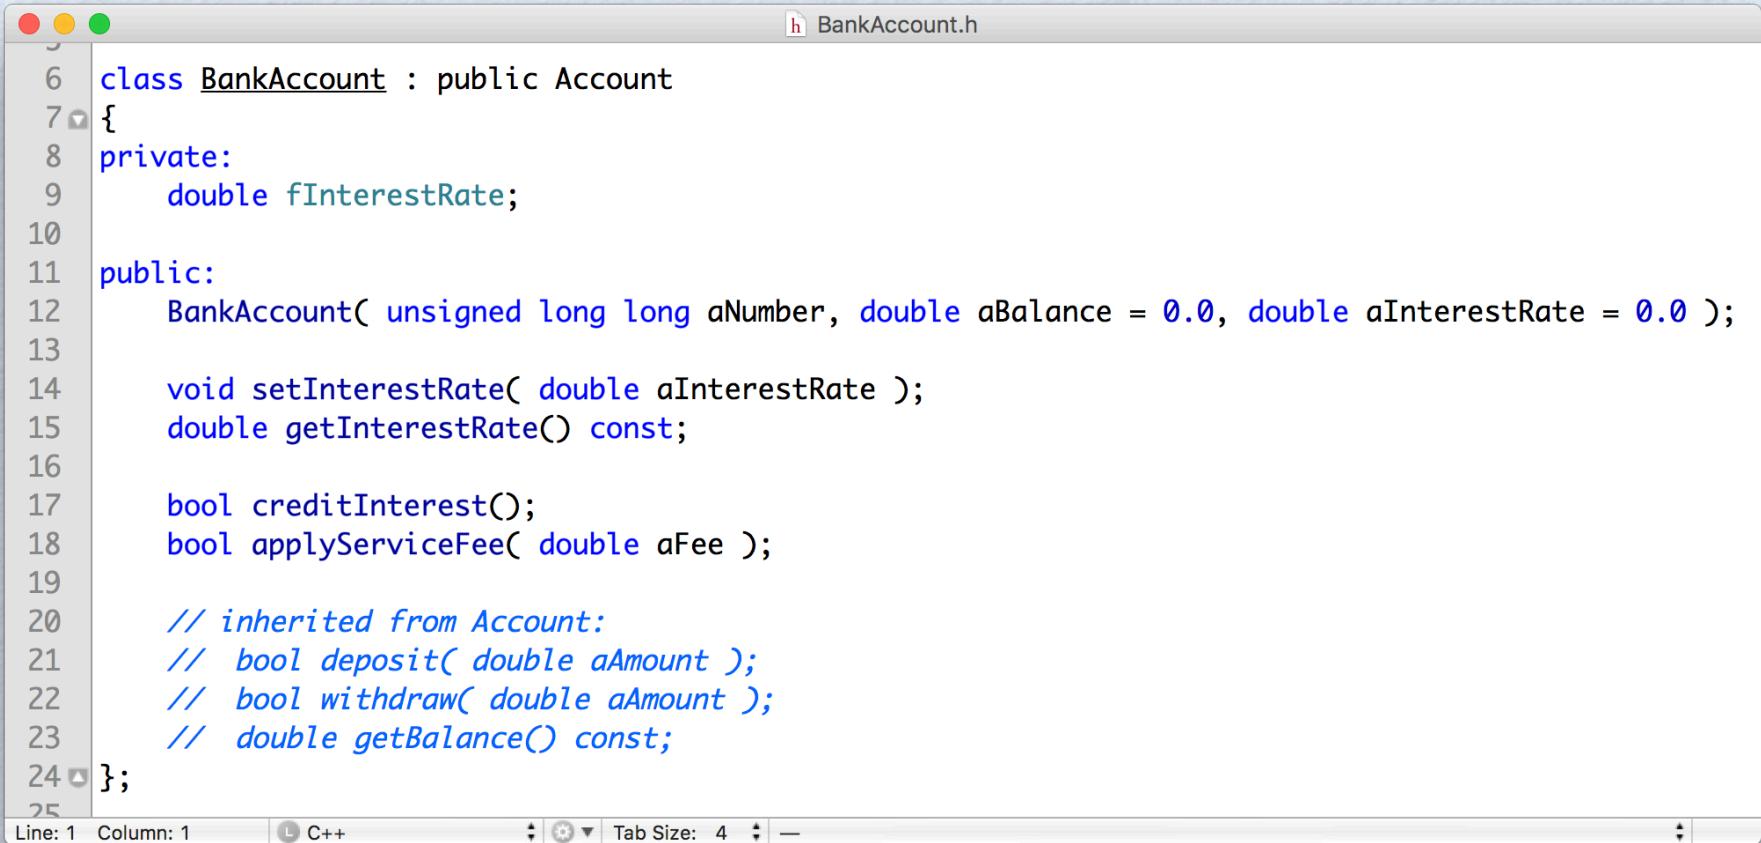
```
4 Account::Account( unsigned long long aNumber, double aBalance )
5 {
6     fNumber = aNumber;
7     fBalance = aBalance;
8 }
9
10 bool Account::deposit( double aAmount )
11 {
12     fBalance += aAmount;
13
14     return true;
15 }
16
17 bool Account::withdraw( double aAmount )
18 {
19     fBalance -= aAmount;
20
21     return true;
22 }
23
24 double Account::getBalance() const
25 {
26     return fBalance;
27 }
28
```

Default in class specification

Signal success

Balance can become negative

A BankAccount Class



The screenshot shows a window titled "BankAccount.h" containing C++ code. The code defines a class `BankAccount` that inherits from `Account`. It includes private members `fInterestRate` and public methods for setting and getting the interest rate, as well as methods for crediting interest and applying service fees. It also includes comments indicating inherited methods from `Account`.

```
6 class BankAccount : public Account
7 {
8     private:
9         double fInterestRate;
10
11    public:
12        BankAccount( unsigned long long aNumber, double aBalance = 0.0, double aInterestRate = 0.0 );
13
14        void setInterestRate( double aInterestRate );
15        double getInterestRate() const;
16
17        bool creditInterest();
18        bool applyServiceFee( double aFee );
19
20        // inherited from Account:
21        // bool deposit( double aAmount );
22        // bool withdraw( double aAmount );
23        // double getBalance() const;
24    };
25
```

Line: 1 Column: 1 C++ Tab Size: 4

- A bank account is a **subtype** of account. A bank account **includes** all features of an account and offers bank account specific additional features (**incremental refinement**).

Access Levels for Inheritance

- **public:**

- Public members in the base class remain public.
- Protected members in the base class remain protected.
- Yields a “is a” relationship, that is, a subtype.

- **protected:**

- Public and protected members in the base class are protected in the derived class.
- Yields a “implemented in terms of” relationship, that is, a new type.

- **private:**

- Public and protected members in the base class become private in the derived class.
- Yields a stricter “implemented in terms of” relationship, that is, a new type.

Public Inheritance

- Public inheritance enables inclusion polymorphism: A subclass is a subtype.
- The subtype relationship is a key ingredient in contemporary object-oriented software development.
- An object of a subclass can be used safely anywhere an object of a superclass is expected. Example, we can use a BankAccount in lieu of an Account object.

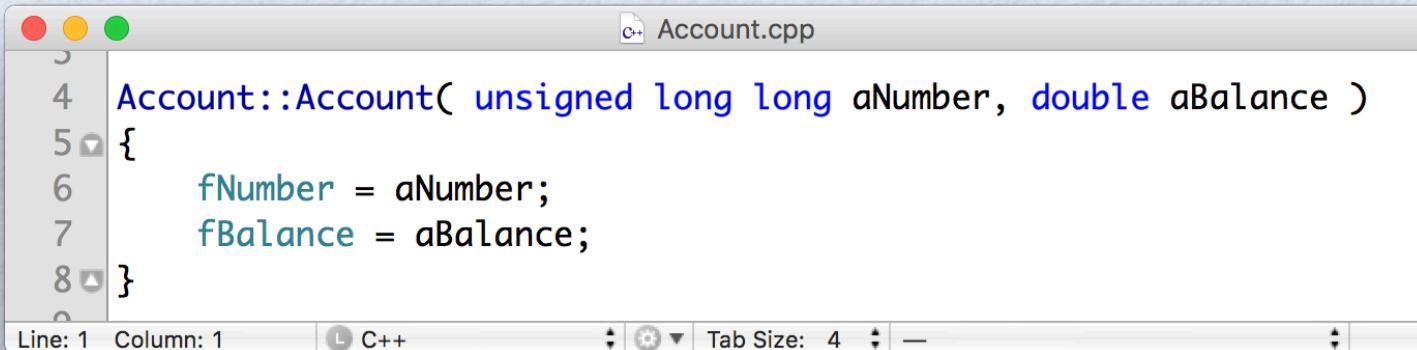
A Grain of Salt

- Public inheritance is not without flaws. It can result in an unwieldy and confusingly large set of public methods.
- Public inheritance is based on additive refinement. New classes are defined by extending an old class with new features.
- Protected and private inheritance allow for a more fine-grained control of the set of public methods. Unfortunately, these mechanisms do not yield subclasses and is often viewed controversial.

Constructors and Inheritance

- Whenever an object of a derived class is instantiated, multiple constructors are called so that each class in the inheritance chain can initialize itself.
- The constructor for each class in the inheritance chain is called beginning with the base class at the top of the inheritance chain and ending with the most recent derived class.

Base Class Initializer

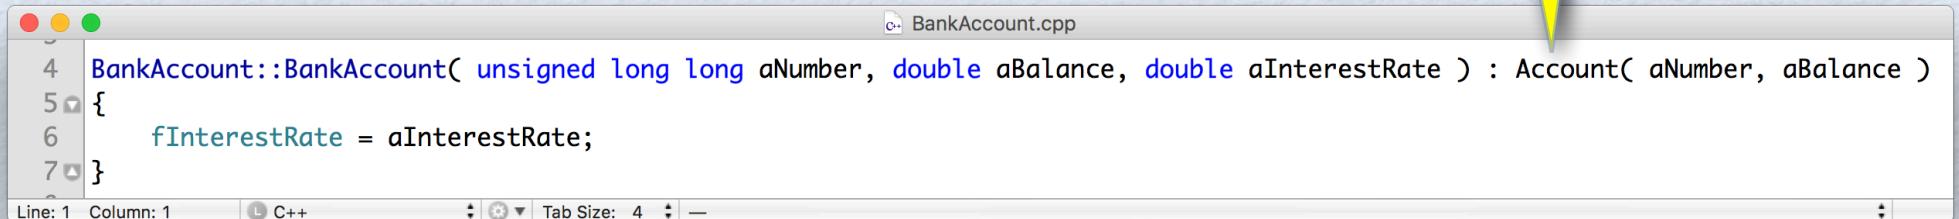


```
4 Account::Account( unsigned long long aNumber, double aBalance )
5 {
6     fNumber = aNumber;
7     fBalance = aBalance;
8 }
```

Line: 1 Column: 1 C++ Tab Size: 4



Direct super class constructor



```
4 BankAccount::BankAccount( unsigned long long aNumber, double aBalance, double aInterestRate ) : Account( aNumber, aBalance )
5 {
6     fInterestRate = aInterestRate;
7 }
```

Line: 1 Column: 1 C++ Tab Size: 4

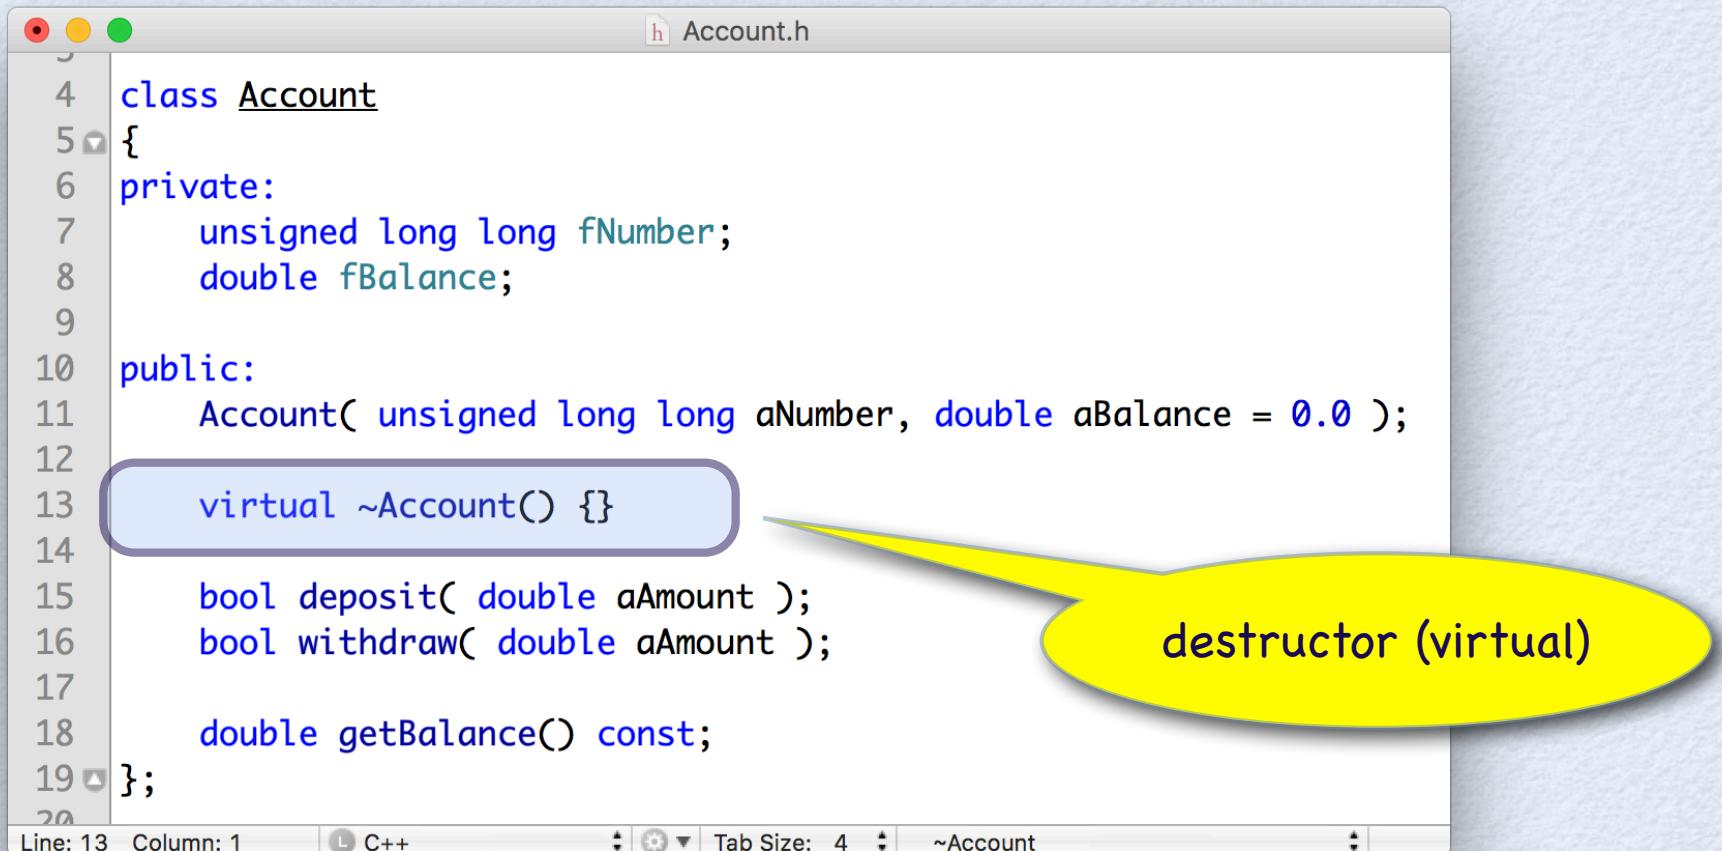
Facts About Base Class Initializers

- If a base class does not have a default constructor, the derived class must provide a base class initializer for it.
- Base class initializers frequently appear alongside member initializers, which use similar syntax.
- If more than one argument is required by a base class constructor, the arguments are separated by comma.
- Reference members need to be initialized using a member initializer.

Destructors and Inheritance

- Whenever an object of a derived class is destroyed, the destructor for each class in the inheritance chain, if defined, is called.
- The destructor for each class in the inheritance chain is called beginning with the most recent derived class and ending with the base class at the top of the inheritance chain.

Account & BankAccount Destructors

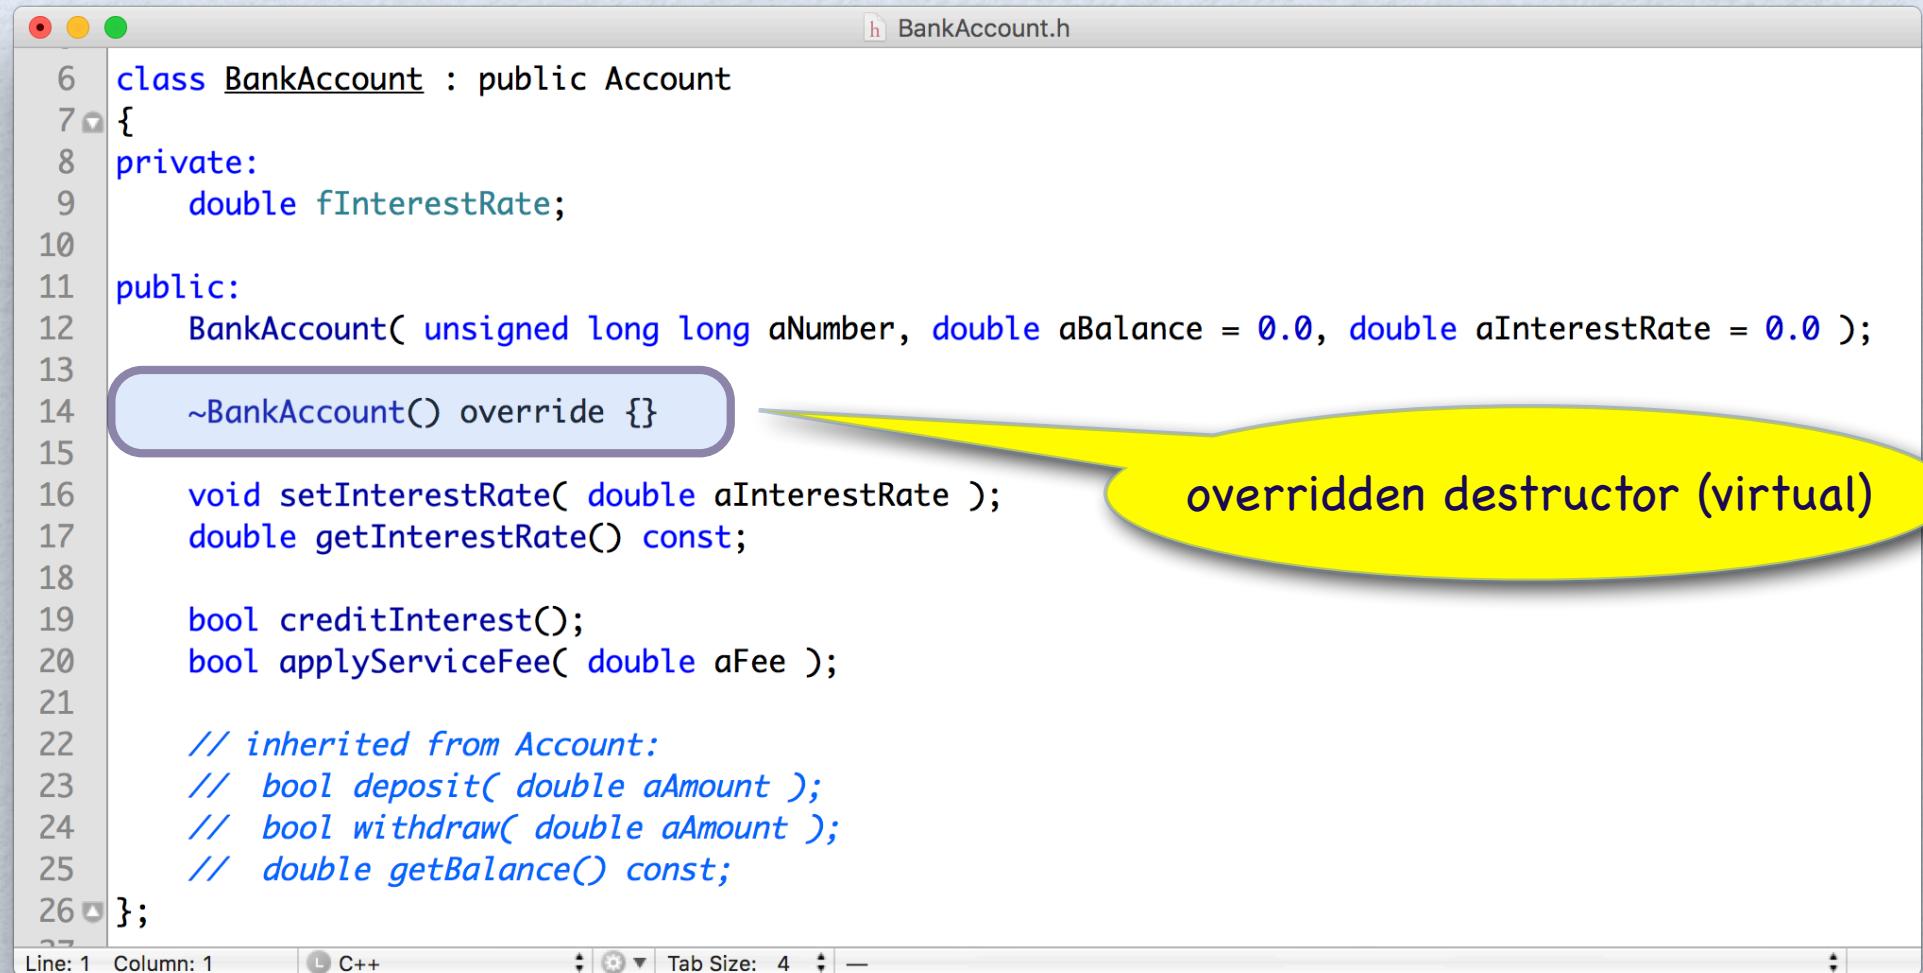


```
4 class Account
5 {
6     private:
7         unsigned long long fNumber;
8         double fBalance;
9
10    public:
11        Account( unsigned long long aNumber, double aBalance = 0.0 );
12
13        virtual ~Account() {}
14
15        bool deposit( double aAmount );
16        bool withdraw( double aAmount );
17
18        double getBalance() const;
19
20};
```

Line: 13 Column: 1 C++ Tab Size: 4 ~Account

A callout bubble with a yellow border and blue text points from the bottom right towards the line `virtual ~Account() {}`. The text inside the bubble is **destructor (virtual)**.

Account & BankAccount Constructors



```
6 class BankAccount : public Account
7 {
8     private:
9         double fInterestRate;
10
11    public:
12        BankAccount( unsigned long long aNumber, double aBalance = 0.0, double aInterestRate = 0.0 );
13
14        ~BankAccount() override {}
15
16        void setInterestRate( double aInterestRate );
17        double getInterestRate() const;
18
19        bool creditInterest();
20        bool applyServiceFee( double aFee );
21
22        // inherited from Account:
23        // bool deposit( double aAmount );
24        // bool withdraw( double aAmount );
25        // double getBalance() const;
26    };
27
```

Line: 1 Column: 1 C++ Tab Size: 4

A yellow callout bubble points to the line `~BankAccount() override {}` with the text "overridden destructor (virtual)".

Virtual Member Functions

- To give a member function from a base class new behavior in a derived class, one overrides it.
- To allow a member function in a base class to be overridden, one must declare the member function `virtual`.
- Note, non-virtual member functions are resolved with respect to the declared type of the object.
- In Java all member functions are virtual.
- Explicitly defining a (virtual) destructor affects which special member functions the compiler synthesizes (C++-11).

The C++11 `override` Keyword

- With C++11, we have a new keyword, `override`, to signal overridden virtual methods (including destructors).
- It has the same semantics as in C#, but you do not need to explicitly specify it.
- The `virtual` specifier can still be used in these cases, but `override` is the new standard.

Virtual Destructors

- When deleting an object using a base class pointer of reference, it is essential that the destructors for each of the classes in the inheritance chain get a chance to run:

```
BankAccount *BAptr;
```

```
Account *Aprt;
```

```
BAptr = new BankAccount( 2.25 );
```

```
Aprt = BAptra;
```

```
...
```

```
delete Aprt;
```

Virtual Account Destructor

```
class Account  
{  
public:  
    virtual ~Account() { ... }  
};
```

Not declaring a destructor virtual is a common source of memory leaks in C++ code.

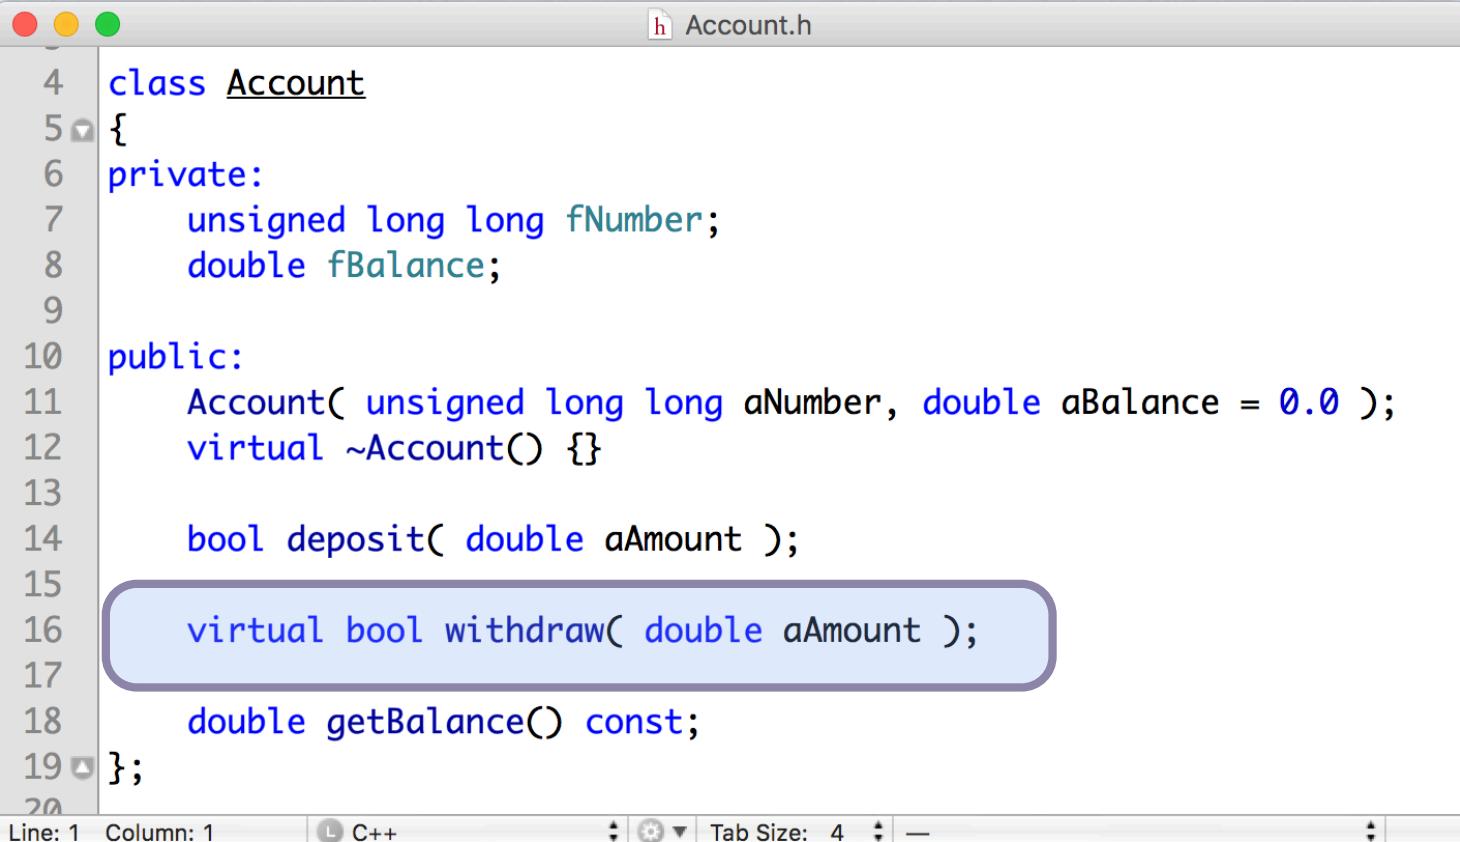
Method Overriding

- Method overriding is an object-oriented programming mechanism that allows a subclass to provide a more specific implementation of a method, which is also present in one of its superclasses.
- The implementation in the subclass overrides (replaces) the implementation in the superclass by providing a method that has the same name, parameter signature, and return type as the method in the parent class. We say these methods belong to the same method family.
- Which (overridden) method is selected at runtime is determined by the receiver object used to invoke it. In general, the most recent definition is chosen, if possible.

Method Family

- A member function of a class always belongs to a specific set, called **method family**. If the elements of this set are **virtual**, then their invocation is governed by **dynamic binding**, a technique that makes polymorphism real.

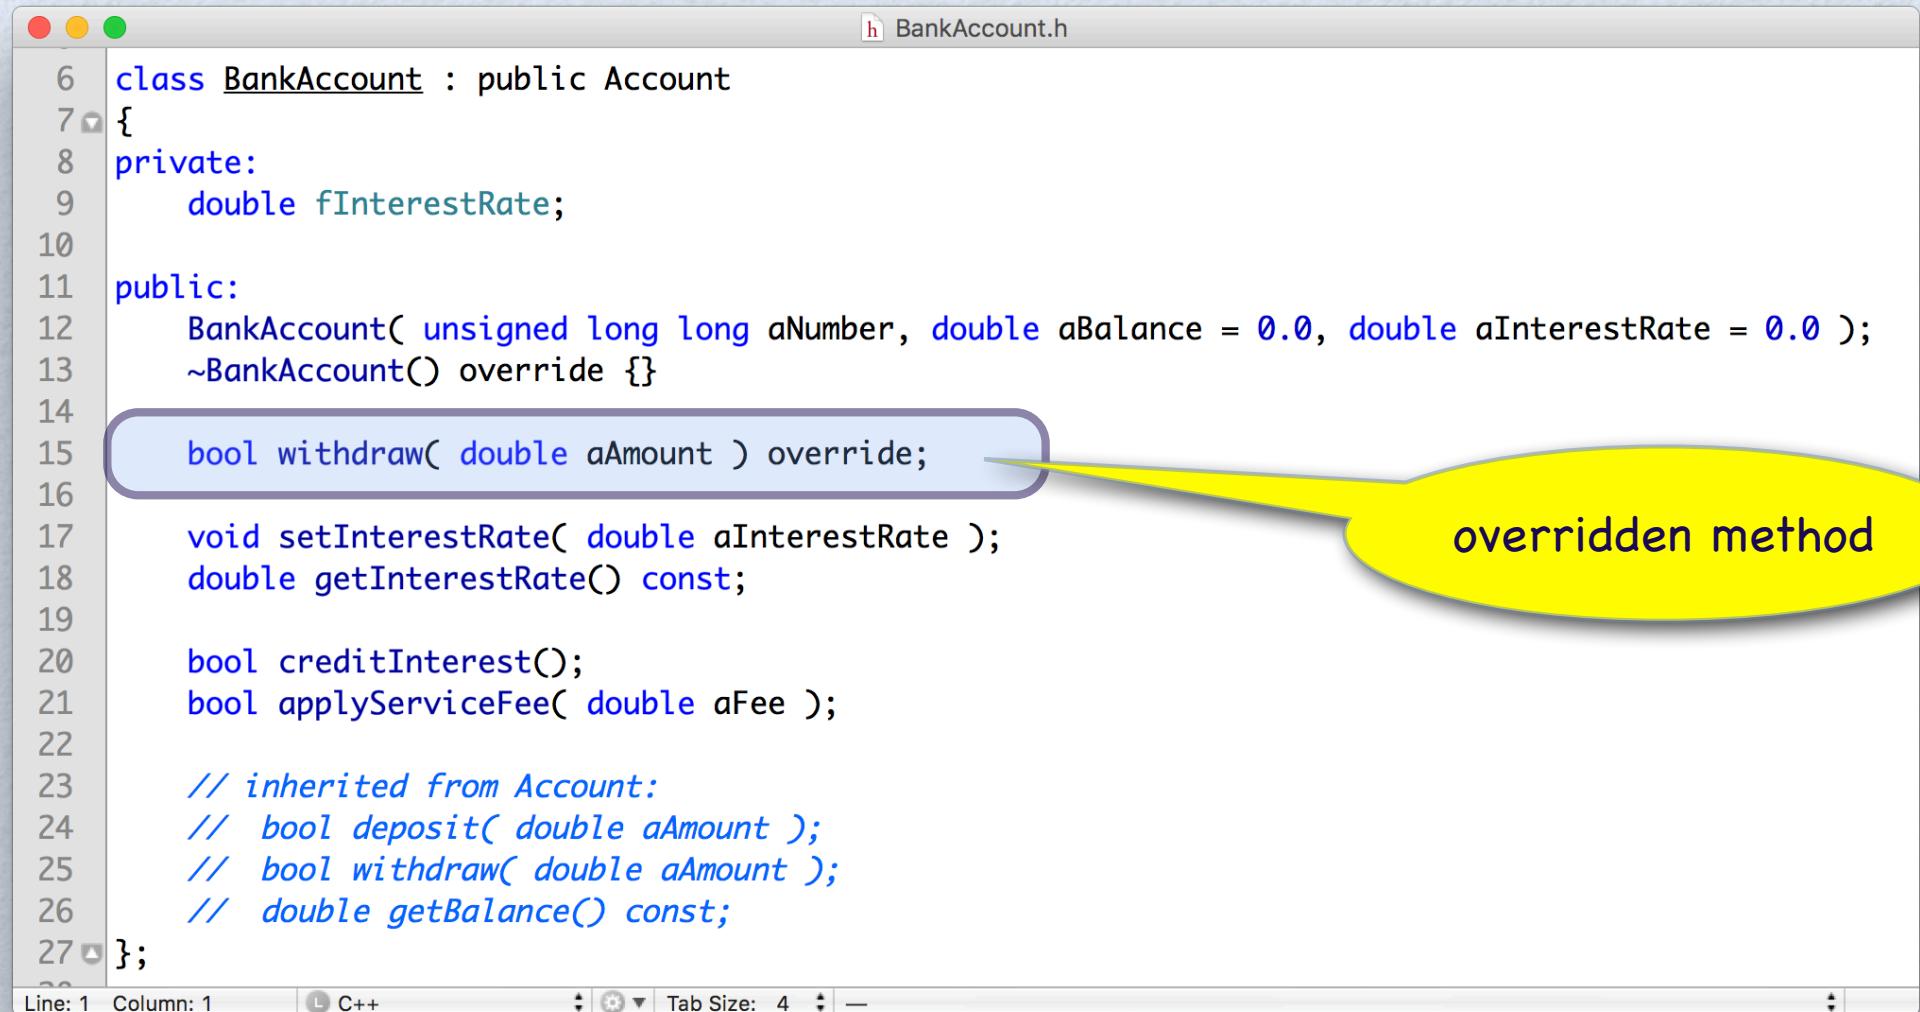
Virtual withdraw Method



```
4 class Account
5 {
6     private:
7         unsigned long long fNumber;
8         double fBalance;
9
10    public:
11        Account( unsigned long long aNumber, double aBalance = 0.0 );
12        virtual ~Account() {}
13
14        bool deposit( double aAmount );
15
16        virtual bool withdraw( double aAmount );
17
18        double getBalance() const;
19    };
20
```

The code editor window shows the `Account.h` header file. The `withdraw` method is highlighted with a purple rounded rectangle. The code defines a class `Account` with private members `fNumber` and `fBalance`, a constructor taking `aNumber` and `aBalance` (with a default value of `0.0`), a destructor, a `deposit` method, a virtual `withdraw` method, and a `getBalance` method marked as `const`. Line numbers are visible on the left, and the status bar at the bottom indicates "Line: 1 Column: 1 C++".

Virtual withdraw Method



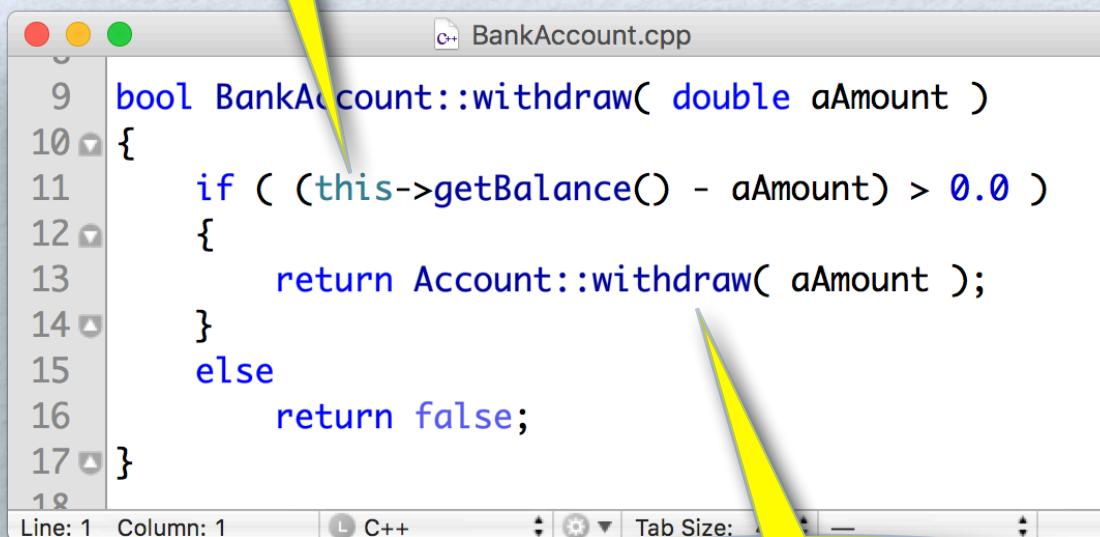
```
6 class BankAccount : public Account
7 {
8     private:
9         double fInterestRate;
10
11    public:
12        BankAccount( unsigned long long aNumber, double aBalance = 0.0, double aInterestRate = 0.0 );
13        ~BankAccount() override {}
14
15        bool withdraw( double aAmount ) override; // overridden method
16
17        void setInterestRate( double aInterestRate );
18        double getInterestRate() const;
19
20        bool creditInterest();
21        bool applyServiceFee( double aFee );
22
23        // inherited from Account:
24        // bool deposit( double aAmount );
25        // bool withdraw( double aAmount );
26        // double getBalance() const;
27};
```

The withdraw Specialization

- In class Account, the method withdraw implements the behavior required to take funds out of an account. The Account class does not do any range checks. Hence, the execution of the withdraw method could result in a negative account balance.
- In general, bank accounts do not allow for negative balances (unless it is a credit account). Hence, a withdrawal can only succeed if there are sufficient funds in the account. The BankAccount's withdraw methods must guard a bank account against requests exceeding the current account balance.

Overriding the withdraw Method

Call inherited method (`this->` optional)



```
9  bool BankAccount::withdraw( double aAmount )
10 { 
11     if ( (this->getBalance() - aAmount) > 0.0 )
12     {
13         return Account::withdraw( aAmount );
14     }
15     else
16         return false;
17 }
18 
```

Call overridden method

Calling a Virtual Method

The screenshot shows a C++ IDE interface with a code editor and a terminal window.

Main.cpp:

```
8 int main()
9 {
10    BankAccount lAccount( 12345, 0.0, 0.5 );
11
12    cout << "Balance: " << lAccount.getBalance() << endl;
13
14    // dynamic method invocation
15
16    Account& lBankAccountReference = lAccount;
17
18    if ( lBankAccountReference.withdraw( 50.0 ) )
19    {
20        cout << "We got instant credit. Wow!" << endl;
21    }
22    else
23    {
24        cout << "The bank refused to give us money." << endl;
25    }
26
27
28    return 0;
29 }
```

A yellow callout bubble points from the line `lBankAccountReference.withdraw(50.0)` to the text "Calls BankAccount::withdraw".

Terminal Output:

```
Kamala:Debug Markus$ ./Inheritance
Balance: 0
The bank refused to give us money.
Kamala:Debug Markus$ _
```

© Dr Markus Lumpe, 2024

Facts About Virtual Members

- Constructors cannot be virtual.
- Declaring a member function `virtual` does not require that this function must be overridden in derived classes, except the member function is declared `pure virtual`.
- Once a member function has been declared `virtual`, it remains `virtual`.
- Parameter and result types must match to properly override a `virtual` member function.
- If one declares a non-`virtual` member function `virtual` in a derived class, the new member function hides the inherited member function.

Note on Virtual Member Functions

- If a virtual member function is called from inside a constructor or destructor, then the version that is run is the one defined for the type of the constructor or destructor itself - static method invocation.
- C++ is a very complex language.